



DEPARTMENT OF INFORMATION ENGINEERING AND COMPUTER
SCIENCE

SOFTWARED AND VIRTUALIZED MOBILE NETWORKS

~ . ~

ACADEMIC YEAR 2021–2022

Morphing Slices Documentation

Professor
Granelli Fabrizio

Student
Vinci Nicolò
220229

July 16, 2022

Contents

Introduction	1
1 Getting Started	3
1.1 Steps for Windows	5
2 Development Workflow	7
2.1 Git and GitHub	7
2.2 Comentsemu Development	10
3 New Topology Building System	13
3.1 src/topology.py	13
3.2 src/topology/topologyParser.py	14
3.3 src/topology/topologyBuilder.py	17
3.4 Build a network topology	20
4 Tested Scenarios	21
4.1 Tested Service	21
4.2 How Migration Works	23
4.3 Scenario 1	24
4.3.1 Running the Scenario	27
4.4 Scenario 2	32
4.4.1 Running the Scenario	33
4.5 Scenario 3	35
4.5.1 Running the Scenario	39
5 Known Issues	43
Conclusions	44
Bibliography	45
List of Figures	48

Introduction

The document provides a detailed description of the Softwarized and Virtualized Mobile Networks project called Morphing Slices. The document is organized as followed:

1. Introduction: it defines a general overview of the theory, the project purpose, the document organization and why the document has been written.
2. Getting Started: it describes how to start the project and which operative system are supported by the project and provides the necessary steps for running the project in Windows.
3. Development Workflow: it reports the workflow followed during the development phase of the project.
4. New Topology Building System: it outlines in details the new innovative way to deploy networks in Comnetsemu.
5. Tested Scenarios: it describes the three tested scenarios and their corresponding results.
6. Known Issues: it reports the issues encountered during the developing.
7. Conclusions: it sums up the entire document focusing on the obtained results and gives a consideration on the new way to deploy networks.

The project is developed on top of Comnetsemu [3] which is a virtual testbed and network emulator for Network Function Virtualization and Software Defined Network. In turn, Comnetsemu extends the network emulator Mininet [12] and allows to deploy Docker container [6] inside its hosts. Docker allows a lightweight virtualization starting container in an isolated environment. A similar approach is taken by Containernet [4] which is a Mininet fork. However, Comnetsemu is focused more on adding essential features to Mininet for a better emulation. To read more about the comparison between Comnetsemu and Mininet: <https://git.comnets.net/public-repo/comnetsemu/-/blob/master/doc/comparison.md>.

<https://git.comnets.net/public-repo/comnetsemu/-/blob/master/doc/faq.md>.

So, the project is strictly related to two paradigms:

- Software Defined Network (SDN): it defines the separation between the Data plane and Control plane. In the networking world, the Data plane involves any activity that results in sending a packet from the end user. Instead, the Control plane handles the Data plane activities without involving end user packets. For example, the Control plane is in charge to make the routing table and the Data plane has to forward the packets following it.
- Network Function Virtualization (NFV): it defines how to virtualize node network functions into building blocks that may be connected, or chained, together to create communication services.

The Network Slicing concept is also implemented in the project and it uses the SDN and NFV paradigms. Network slicing means splitting the network in different slices providing logical end-to-end connection. It allows to deploy multiple logical, self-contained networks on top a common hardware. It supports flexible on-demand networks and isolation among them. Within a slice, the forwarding rules can be defined with flows. Then, they may be manipulated dynamically to redirect the traffic. The OpenFlow protocol [13] was created to deploy the Network Slicing concept. So, the figure 1 summarizes the theoretical idea behind the project.




Figure 1: Introduction summary

The goal of the project is to enable the RYU SDN controller[18] to build network slices and dynamically modify the network traffic. The SDN controller will not only slice, but reprogram connectivity within the slice. Considering to have the same service deployed in two different servers, a service migration can be performed manipulating and reprogramming the flows within the slice. So, the project aims at demonstrating a dynamic, stateful and transparent service migration:

- Dynamic: migration is performed after a user input.
- Stateful: service state is kept after the migration.
- Transparent: client is not aware of the migration.

It has been decided to write the document for giving a detailed documentation about the project. By the way, reading the document it can be useful for the future Comnetsemu project developers. A very smooth development workflow and new convenient method of developing networks have been described. So, future developers may be inspired by the document during their projects.

Chapter 1

Getting Started

The project can be run on Linux-based and Windows operative system. However, some extra steps need to be taken on Windows. They are described in the section 1.1. On Windows, it is recommended to first reading the section 1.1 and then follow the steps below. On a Linux-based, the below steps can be performed immediately. Firstly, it is necessary to clone Github repository:

```
$ git clone --recursive https://github.com/nico989/SVMN.git
$ cd SVMN
```

The option *recursive* is to clone recursively any git submodule in a repository. The project imports as a git submodule the original Comnetsemu Github repository from:

<https://github.com/stev lorenz/comnetsemu>.

Then, the folder *scripts* needs the execution permissions:

```
$ chmod -R +x scripts
```

At the end, the initialization script must be run:

```
$ scripts/init.sh
```

At the first time the *init.sh* is executed, the local project folder might not be copied in the Comnetsemu virtual machine. So, the error showed at figure may be prompted.

```
nico989@ubuntu:~/SVMN$ scripts/init.sh
[2022-05-29 18:19:39.396][INFO] [main:18 ] == Checking tools ==
[2022-05-29 18:19:39.399][INFO] [assert_tool:28 ] 'vagrant' found at /usr/bin/vagrant
[2022-05-29 18:19:39.403][INFO] [assert_tool:28 ] 'python3' found at /usr/bin/python3
[2022-05-29 18:19:39.406][INFO] [assert_tool:28 ] 'pip3' found at /home/nicolo/.local/bin/pip3
[2022-05-29 18:19:39.409][INFO] [assert_tool:28 ] 'pipenv' found at /home/nicolo/.local/bin/pipenv
[2022-05-29 18:19:39.411][INFO] [assert_tool:28 ] 'shellcheck' found at /usr/bin/shellcheck
[2022-05-29 18:19:39.414][INFO] [assert_tool:28 ] 'inotifywait' found at /usr/bin/inotifywait
[2022-05-29 18:19:39.416][INFO] [assert_tool:28 ] 'rsync' found at /usr/bin/rsync
[2022-05-29 18:19:39.418][INFO] [assert_tool:28 ] 'gen' found at /usr/bin/gen
[2022-05-29 18:19:39.422][INFO] [main:25 ] == Structure ==
[2022-05-29 18:19:39.424][WARN] [main:30 ] '.venv' directory already exists
[2022-05-29 18:19:39.426][INFO] [main:34 ] == Python packages ==
[2022-05-29 18:19:39.427][INFO] [main:35 ] Installing Python packages
Installing dependencies from Pipfile.lock (1f4b2f)...
  0/6 -
To activate this project's virtualenv, run `source shell`.
Alternatively, run a command inside the virtualenv with `pipenv run`.
[2022-05-29 18:19:40.788][INFO] [main:38 ] Installing 'pre-commit'
pre-commit installed at .git/hooks/pre-commit
[2022-05-29 18:19:41.557][INFO] [main:42 ] == comnetsemu ==
[2022-05-29 18:19:41.559][INFO] [main:43 ] Initializing 'comnetsemu'
bringing machine 'comnetsemu' up with 'virtualbox' provider.
==> comnetsemu: Checking if box 'bento/ubuntu-20.04' version '202107.28.0' is up to date...
==> comnetsemu: Machine already provisioned. Run 'vagrant provision' or use the "--provision"
==> comnetsemu: flag to force provisioning. Provisioners marked to run always will still run.

==> comnetsemu: Machine 'comnetsemu' has a post 'vagrant up' message. This is a message
==> comnetsemu: from the creator of the Vagrantfile, and not from Vagrant itself.
==> comnetsemu:
==> comnetsemu: VM already started! Run "$ vagrant ssh comnetsemu" to ssh into the running VM.
==> comnetsemu:
==> comnetsemu: **IMPORTANT!!** For all Comnetsemu users and developers:
==> comnetsemu:
==> comnetsemu: **Please** run the upgrade process described [here](https://git.comnets.net/public-repo/comnetsemu/upgrade-comnetsemu-and-dependencies) when there is a new release
==> comnetsemu: published [here](https://git.comnets.net/public-repo/comnetsemu/~tags).
==> comnetsemu: New features, fixes and other improvements require run the upgrade script **manually**.
==> comnetsemu: But the script will check and perform upgrade automatically and it does not take much time if you have a good network connection.
==> comnetsemu:
sudo: comnetsemu/app/morphing_slices/scripts/init.sh: command not found
Connection to 172.19.0.11 closed.
[2022-05-29 18:19:46.016][INFO] [main:45 ] Error initializing 'comnetsemu'
```

Figure 1.1: First *init.sh* run

To copy the local project folder in Comnetsemu the first time, the *dev.sh* script must be run before the *init.sh* script:

```
$ scripts/dev.sh
```

Now, the *init.sh* script can finish correctly. Read the chapter 2 for more details about the *dev.sh* script. The *init.sh* script performs different preliminary operations:

1. It checks if the required packages are installed.

```
13 # Assert tool(s)
14 INFO "=== Checking tools ==="
15 assert_tool vagrant
16 assert_tool python3
17 assert_tool pip3
18 assert_tool pipenv
19 assert_tool shellcheck
20 assert_tool inotifywait
21 assert_tool rsync
22 assert_tool gem
```

Figure 1.2: Packages required

2. Then, it creates the Python virtual environment folder and install the Python packages inside it. It also installs the pre-commit package, check the chapter 2 for more information.

```
24 # Create .venv directory
25 INFO "=== Structure ==="
26 if [ ! -d ".venv_dir" ]; then
27     INFO "creating '.venv' directory"
28     mkdir "${VENV_DIR}"
29 else
30     WARN "'.venv' directory already exists"
31 fi
32
33 # Install Python packages
34 INFO "=== Python packages ==="
35 INFO "Installing Python packages"
36 (pipenv install -dev) || { FATAL "Error installing python packages"; exit 1; }
37 # install pre-commit
38 INFO "Installing 'pre-commit'"
39 (pipenv run pre-commit install) || { FATAL "Error installing 'pre-commit'"; exit 1; }
```

Figure 1.3: Set Python virtual environment

3. At the end, it starts the Comnetsemu virtual machine and run the initialization script for it.

```
41 # Create and configure comnetsemu
42 INFO "=== comnetsemu ==="
43 INFO "Initializing 'comnetsemu'"
44 (cd "${_DIRNAME}/../comnetsemu" && vagrant up) || { FATAL "Error initializing 'comnetsemu'; exit 1; }
45 (cd "${_DIRNAME}/../comnetsemu" && vagrant ssh -t "sudo comnetsemu/app/morphing_slices/scripts/init.sh") || { FATAL "Error initializing 'comnetsemu'; exit 1; }
```

Figure 1.4: Run and configure Comnetsemu

So, there is another installation script located at *src/scripts/init.sh*. The script performs four crucial steps for the Comnetsemu virtual machine:

1. It updates the Comnetsemu machine and installs the *parallel* package.

```
14 # Packages
15 INFO "Installing packages"
16 sudo apt-get update \
17 && sudo apt-get install -y \
18 parallel
```

Figure 1.5: Update and install Comnetsemu packages

2. If they are not present yet, it builds up the Docker Images for the Mininet hosts.

```

21 # check docker images
22 if [[ "$(docker images -q ${DEV_HOST_IMAGE} 2>/dev/null)" == "" ]]; then
23     WARN "docker image '${DEV_HOST_IMAGE}' does not exists"
24     INFO "Building Docker image '${DEV_HOST_IMAGE}'"
25     docker build --rm -t ${DEV_HOST_IMAGE} -f "${DIRNAME}/../docker/Dockerfile_dev_host" .
26 fi
27 if [[ "$(docker images -q ${DEV_SERVER_IMAGE} 2>/dev/null)" == "" ]]; then
28     WARN "docker image '${DEV_SERVER_IMAGE}' does not exists"
29     INFO "Building Docker image '${DEV_SERVER_IMAGE}'"
30     docker build --rm -t ${DEV_SERVER_IMAGE} -f "${DIRNAME}/../docker/Dockerfile_dev_server" .
31 fi

```

Figure 1.6: Build Docker images

3. It installs the Python packages from the requirements file generated by *scripts/dev.sh* and two more packages.

```

32 INFO "Installing Python dependencies from 'requirements.txt'"
33 sudo pip install -r "${DIRNAME}/../requirements.txt"
34 sudo pip install eventlet==0.30.2 Flask==2.0.3

```

Figure 1.7: Install Python packages in Comnetsemu

4. It patches the RYU Python package. Check the chapter 5 for more details.

```

36 INFO "Checking Ryu"
37 # Ryu gui topology html directory
38 RYU_GUI_TOPOLOGY_HTML_DIR="${dirname "$(python3 -c "import ryu; print(ryu.__file__)")"/app/gui_topology/html}"
39 readonly RYU_GUI_TOPOLOGY_HTML_DIR
40 if [ ! -d "${RYU_GUI_TOPOLOGY_HTML_DIR}" ]; then
41     WARN "Ryu not fixed, fixing..."
42     # html folder
43     sudo mkdir "${RYU_GUI_TOPOLOGY_HTML_DIR}"
44     sudo wget --quiet --show-progress -P "${RYU_GUI_TOPOLOGY_HTML_DIR}" https://raw.githubusercontent.com/faucetsdn/ryu/master/ryu/app/gui_topology/html/index.html
45     sudo wget --quiet --show-progress -P "${RYU_GUI_TOPOLOGY_HTML_DIR}" https://raw.githubusercontent.com/faucetsdn/ryu/master/ryu/app/gui_topology/html/router.svg
46     sudo wget --quiet --show-progress -P "${RYU_GUI_TOPOLOGY_HTML_DIR}" https://raw.githubusercontent.com/faucetsdn/ryu/master/ryu/app/gui_topology/html/ryu_topology.css
47     sudo wget --quiet --show-progress -P "${RYU_GUI_TOPOLOGY_HTML_DIR}" https://raw.githubusercontent.com/faucetsdn/ryu/master/ryu/app/gui_topology/html/ryu_topology.js
48     INFO "Ryu fixed"
49 else
50     INFO "Ryu already fixed"
51 fi

```

Figure 1.8: Fix RYU Python package

1.1 Steps for Windows

First of all, the user must install Windows Subsystem for Linux (WSL) version 2 following the Microsoft official guide:

<https://docs.microsoft.com/it-it/windows/wsl/install>

Then, the next configurations must be done in the WSL and not in Windows directly:

1. Install Vagrant in WSL:

<https://www.vagrantup.com/downloads>.

It can be installed from the Ubuntu repositories as well, but the version may be older than the latest. So, it is recommended to download the zip file and install the deb package manually. Vagrant must be installed in Windows as well.

2. Configure Vagrant to run in WSL following the official guide:

<https://www.vagrantup.com/docs/other/wsl>.

Two environment variable must be set for the basics configuration:

```

$ export VAGRANT_WSL_ENABLE_WINDOWS_ACCESS="1"
$ export PATH="$PATH:/mnt/c/Program Files/Oracle/VirtualBox"

```

3. Install the Vagrant plugin for VirtualBox in the WSL:

https://github.com/Karandash8/virtualbox_WSL2.

The following command can be run to install the plugin:

```
$ vagrant plugin install virtualbox_WSL2
```

4. The line 565 of *platform.rb* file located at
*/opt/vagrant/embedded/gems/[VAGRANT_VERSION]/gems/
vagrant-[VAGRANT_VERSION]/lib/vagrant/util/platform.rb*
must be replaced from:

```
if info && (info[:type] == "drvfs" || info[:type] == "9p")
```

To:

```
if info && (info[:type] == "drvfs" || info[:type] == "9p"  
|| info[:type] == "ext4")
```

More information about the issue:

<https://github.com/hashicorp/vagrant/issues/11623>

After completing all the steps, the repository can be cloned and scripts can be run as described above.

Chapter 2

Development Workflow

The chapter describes the workflow followed during the project development. The purpose is to suggest a smooth and comfortable way of working in the Comnetsemu environment. Then, it is up to the developers to follow the proposed approach or create a new one.

2.1 Git and GitHub

Git has been used for versioning the code and the repository is hosted on GitHub. Git allows to keep tracking and managing the code. GitHub is useful to host on the web the repository and integrate different workflows [9]. A GitHub workflow is a configurable automated process that run one or more jobs and it is defined by a YAML file. The project integrates two workflows useful for testing and deployment purpose:

- Continuous Integration (CI) [2]: it is defined by the *ci.yml* file located in the *.github/workflows* folder. CI workflow automatically builds and tests the software every time a developer pushes changes to the main branch. So, every time the workflow performs the steps defined in the YAML file after a push action or pull request. The CI mainly tests the installation of packages for this project.

```
15 steps:
16   - name: Checkout repository
17     uses: actions/checkout@v2
18     with:
19       submodules: recursive
20
21   - name: Install packages
22     run: |
23       sudo apt-get update
24       sudo apt-get install ruby-dev shellcheck
25       shell: bash
26
27   - name: Install pipenv
28     run: pipx install pipenv
29
30   - name: Python
31     uses: actions/setup-python@v2
32     with:
33       python-version: "3.8"
34       cache: "pipenv"
35
36   - name: Install Python packages
37     run: pipenv install --dev
38
39   - name: pre-commit
40     run: pipenv run pre-commit run --all-files
```

Figure 2.1: CI steps

The result can be checked on the GitHub user interface under the tab *Action* and the item *ci*.




Figure 2.2: CI result

If there is a red cross like in the figure 2.2, it means that something went wrong. Otherwise, if all went good, there will be a green tick.

- Continuous Deployment (CD) [1]: it is defined by the *release.yml* file located in the *.github/-workflows* folder. It generates *morphing_slices.tar.gz* file for any new tagged version. So, the workflow is triggered every time someone performs a push action with the option *-v* for defining the version.

```
$ git tag v0.0.30
$ git push --origin tags
```

After the pushing request, the steps reported in the figure 2.3 are taken.

```
15 steps:
16   - name: Checkout repository
17     uses: actions/checkout@v2
18     with:
19       submodules: recursive
20
21   - name: Environment
22     run: echo "RELEASE_VERSION=${GITHUB_REF#refs/*/}" >> $GITHUB_ENV
23
24   - name: Install pipenv
25     run: pipx install pipenv
26
27   - name: Python
28     uses: actions/setup-python@v2
29     with:
30       python-version: "3.8"
31       cache: "pipenv"
32
33   - name: Scripts permissions
34     run: chmod -R +x scripts
35
36   - name: Generate morphing_slices.tar.gz
37     run: scripts/prod.sh
38
39   - name: Release
40     uses: softprops/action-gh-release@v1
41     with:
42       body: Morphing Slices ${ env.RELEASE_VERSION }
43       files: |
44         morphing_slices.tar.gz
```

Figure 2.3: CD steps

The most important is the one highlighted in red. It executes the script *prod.sh* located in the folder *scripts*. The script core function is to generate the *.tar.gz* file from the *src* folder.

```
44 # Generate final tar|gz file
45 INFO "generating  ${TAR_GZ_NAME}"
46 find "${TMP_DIR}" -printf "%p\n" | tar -czf "${TAR_GZ_NAME}" --no-recursion -C "${TMP_DIR}" -T -
```

Figure 2.4: Folder generation

The result can be checked on the GitHub user interface under the tab *Action* and the item *release*.

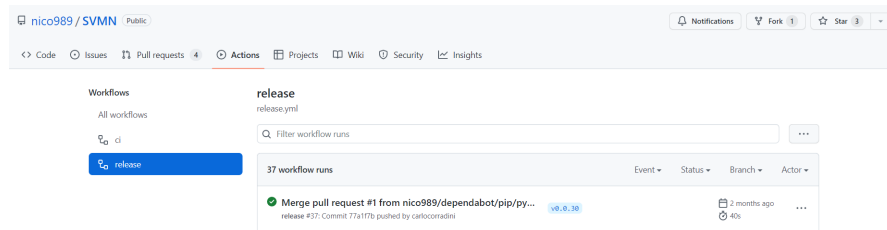


Figure 2.5: CD result

If there is a red cross like in the figure 2.2, it means that something went wrong. Otherwise, if the *.tar.gz* file has been created correctly, there will be a green tick as reported in figure 2.5. So, any project version can be easily downloaded from the GitHub repository.

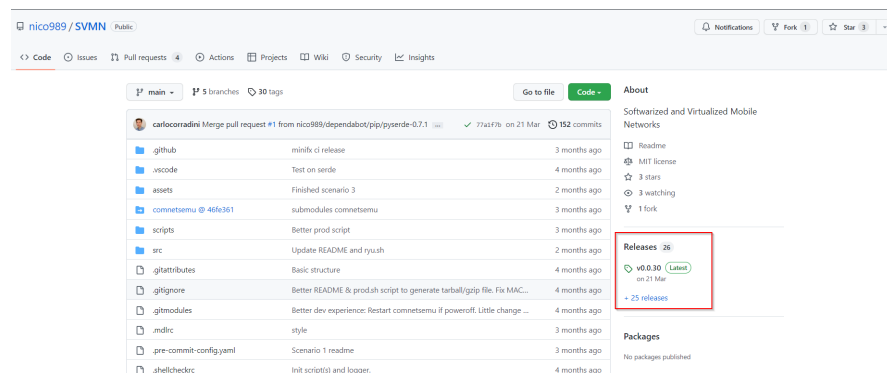


Figure 2.6: CD release

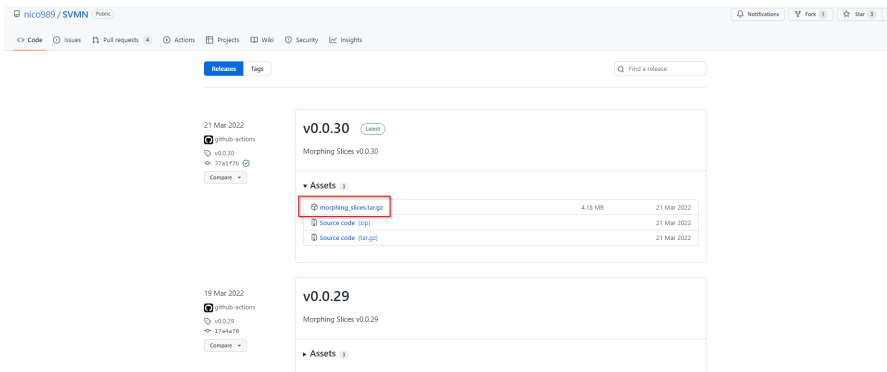


Figure 2.7: CD release download

Then, other two tools are used to maintain the code and the project itself:

- Pre-commit [15]: it is a framework to manage and maintain multi-language Pre-commit hooks. A Git hook script is useful to identify simple issues before submitting the code. Hooks are defined in the *pre-commit-config.yaml* file. They run automatically on every commits pointing out issues such as missing semicolons, trailing whitespace, and debug statements.

```

3 repos:
4   - repo: https://github.com/pre-commit/pre-commit-hooks
5     rev: v4.1.0
6     hooks:
7       - id: trailing-whitespace
8       - id: end-of-file-fixer
9       - id: check-docstring-first
10      - id: check-yaml
11      - id: debug-statements
12      - id: name-tests-test
13      - id: requirements-txt-fixer
14   - repo: https://github.com/jumanjihouse/pre-commit-hooks
15     rev: 2.1.5
16     hooks:
17       - id: script-must-have-extension
18       - id: markdownlint
19       - id: shellcheck
20   - repo: https://github.com/psf/black
21     rev: 22.1.0
22     hooks:
23       - id: black
24   - repo: https://github.com/Lucas-C/pre-commit-hooks-nodejs
25     rev: v1.1.1
26     hooks:
27       - id: dockerfile_lint

```

Figure 2.8: Pre-commit YAML file

- Dependabot [5]: it is defined by the *dependabot.yml* file located in the *.github* folder. It is used to check the third-party libraries updates. It checks the *pip* and *docker* dependencies weekly. A pull request will be created, if an update is found.

```

2 updates:
3   # Enable version updates for pipenv
4   - package-ecosystem: "pip"
5     # Look for a "Pipfile" in the "root" directory
6     directory: "/"
7     # Check for updates once a week
8     schedule:
9       interval: "weekly"
10
11  # Enable version updates for Docker
12  - package-ecosystem: "docker"
13    # Look for a "Dockerfile" in the "root" directory
14    directory: "/"
15    # Check for updates once a week
16    schedule:
17      interval: "weekly"

```

Figure 2.9: Dependabot YAML file

2.2 Comentsemu Development

In this section the development process in the Comnetsemu virtual machine is described. It might be annoying to copy and paste every time the local project in the Comnetsemu virtual machine. However, Comnetsemu has a shared folder called *comnetsemu* already set up. It means that any file in that folder can be used in the virtual machine as well. So, the simple idea is to set up a watcher which automatically copies and pastes any update from the *src* folder in the *comnetsemu* shared folder.




Figure 2.10: Development schema

The script *dev.sh* located in the folder *scripts* must be run to start the development process:

1. First, it checks from the Vagrant status if the Comnetsemu virtual machine has been started. If it is powered off, the script will restart Comnetsemu.

```

20 # Check Vagrant status
21 INFO "=== Checking Vagrant ==="
22 case "${VAGRANT_STATUS}" in
23     running)
24         INFO "'comnetsemu' is running"
25         ;;
26     poweroff)
27         # Restarting comnetsemu
28         INFO "'comnetsemu' is poweroff, restart..."
29         INFO "Restarting 'comnetsemu'"
30         (cd "${_DIRNAME}/../comnetsemu" && vagrant up) || { FATAL "Error restarting 'comnetsemu'"; exit 1; }
31         ;;
32     *)
33         ERROR "Vagrant status '${VAGRANT_STATUS}' not checked"
34         exit 1
35         ;;
36 esac

```

Figure 2.11: Check Comnetsemu status

2. It generates the *requirements.txt* file which contains the required packages to be installed.

```

38 # Generate requirements.txt
39 INFO "Generating 'requirements.txt' from 'pipenv'"
40 pipenv lock -r > "${_DIRNAME}/../src/requirements.txt"

```

Figure 2.12: Generate requirements

3. Then, it synchronizes the destination directory *morphing_slices* in the shared folder *comnetsemu/app* with the local directory *src*. The first time, the local directory is copied in the shared folder. The utility *rsync* [17] is used for synchronization purpose.

```

42 # Sync watch and destination directories
43 INFO "=== Sync ==="
44 function sync() {
45     DEBUG "Syncing '$WATCH_DIR' with destination '$DEST_DIR'"
46     rsync --archive --verbose --compress --delete --human-readable --quiet "$WATCH_DIR" "$DEST_DIR"
47 }
48 # Sync directory
49 INFO "Syncing directory '$WATCH_DIR'"
50 sync
51

```

Figure 2.13: Rsync synchronization

4. It initializes Comnetsemu running the *init.sh* script located in the *scripts* folder.

```

53 # Initialize comnetsemu
54 INFO "Initializing 'comnetsemu'"
55 (cd "${_DIRNAME}/../comnetsemu" && vagrant ssh -- -t 'cd comnetsemu/app/morphing_slices && sudo scripts/init.sh') || { FATAL "Error initializing 'comnetsemu'"; exit 1; }

```

Figure 2.14: Initialize Comnetsemu

5. At the end, the watcher starts to monitor the *src* local folder. The command *inotifywait* [10] is used to monitor any change in *src*. Every time a change occurs, only the change will be reflected in the destination folder by *rsync*.

```

57 # Watcher
58 INFO "Starting watcher on '$WATCH_DIR' with destination '$DEST_DIR'"
59 while inotifywait --recursive --quiet --event modify,create,delete,move "$WATCH_DIR"; do
60     sync
61 done
62 INFO "Watcher stopped";

```

Figure 2.15: Start watcher

When the *dev.sh* script runs, the output looks like the one showed in the figure 2.16 at the end.

```
[2022-05-29 18:17:24.677][INFO ][main:58 ] Starting watcher on '/home/nicolo/SVMN/src/' with destination '/home/nicolo/SVMN/comnetsemu/app/morphing_slices/'
```

Figure 2.16: Run *dev.sh*

If a change occurs, it will be prompted in the output as shown in the figure 2.17.

```
[2022-05-29 18:17:24.677][INFO ][main:58 ] Starting watcher on '/home/nicolo/SVMN/src/' with destination '/home/nicolo/SVMN/comnetsemu/app/morphing_slices/'  
/home/nicolo/SVMN/src/ CREATE test.txt  
/home/nicolo/SVMN/src/ MODIFY test.txt  
/home/nicolo/SVMN/src/ DELETE test.txt
```

Figure 2.17: Change occurrence

Chapter 3

New Topology Building System

The project brings a new innovative way to deploy virtual networks in the Comnetsemu environment. Comnetsemu basically uses the Mininet python API to create the network topology and emulate the nodes. It is needed to write a long python file to emulate a big network topology and it is not easily readable for future developers or debugging purpose. So, the idea is to take an approach similar to Kubernetes [11] having a YAML or JSON file to describe the network topology to achieve a better readability and avoid python boiler plate code. Therefore, it is more easy to update the topology modifying the YAML/JSON file than a long python code. The idea of the new building system is to parse the YAML/JSON file and create the network using the Comnetsemu python API. So, there are three main python scripts:

- *src/topology.py*.
- *src/topology/topologyParser.py*.
- *src/topology/topologyBuilder.py*.

Hence, a new layer has been added to improve the usability in deploying network topology in the Comnetsemu environment. Future developers will need to write a single YAML/JSON file for the network topology. A practical application of the new topology building system can be seen in the chapter 4,

3.1 *src/topology.py*

This is entry point for building the topology. Firstly, a parser has been defined in order to get the YAML/JSON file as input.

```
30 if __name__ == "__main__":
31     # Arguments
32     parser = argparse.ArgumentParser(
33         description="Topology builder",
34         formatter_class=lambda prog: argparse.HelpFormatter(prog, max_help_position=40),
35     )
36     parser.add_argument(
37         "-f", help="Topology file", required=True, action="store", type=str
38     )
39     args = parser.parse_args()
```

Figure 3.1: Main *topology.py*

The input file is parsed in the *getTopology* function to get a topology object in order to feed the *buildTopology* function to emulate the network. The latter is described in the section 3.3.

```

41 # Topology
42 logger.info(f"Analyzing topology file '{args.file}'")
43 topology = getTopology(args.file)
44 (network, manager) = topologyBuilder.buildTopology(topology)

```

Figure 3.2: Get and build topology

At the end, the `Containernet` object returned by `buildTopology` and Mininet command line interface are started. The `cleanTopology` function is called when the user closes the CLI.

```

46 # Network
47 network.start()
48 cli.CLI(network)
49 topologyBuilder.cleanTopology(topology, network, manager)
50

```

Figure 3.3: Start network and Mininet CLI

Regarding the `getTopology` function, it takes in input the YAML/JSON file, it analyzes it with the `fileFormat` function and calls the `parseTopology` to deserialize the file in the topology object. More details of the `parseTopology` are reported in the section 3.2.

```

20 def getTopology(file: str) -> topologyBuilder.Topology:
21     topology = None
22     format = fileFormat(file)
23
24     with open(file) as f:
25         topology = f.read()
26
27     return topologyParser.parseTopology(topology, format)

```

Figure 3.4: `getTopology` function

To conclude, the `fileFormat` function analyzes the file type provided as input by the user. The accepted extensions are `.yaml`, `.yml` and `.json`, because the `parseTopology` is able to deserialize only them. If a user inputs another file type, an exception is returned.

```

10 def fileFormat(file: str) -> topologyParser.Format:
11     _, file_extension = os.path.splitext(file)
12     if file_extension == ".yaml" or file_extension == ".yml":
13         return topologyParser.Format.YAML
14     elif file_extension == ".json":
15         return topologyParser.Format.JSON
16     else:
17         raise ValueError(f"Unknown file extension {file_extension}")

```

Figure 3.5: `fileFormat` function

3.2 src/topology/topologyParser.py

The script aims to deserialize the input file into the custom topology object. It relies on the `pyserde` [16] python library which allows to deserialize YAML/JSON file into python object and viceversa. The library provides two functions called `from_yaml` and `from_json` to do the deserialization and get the topology object.


```

75 class Format(enum.Enum):
76     YAML = enum.auto()
77     JSON = enum.auto()
78
79
80 def parseTopology(topology: str, format: Format) -> Topology:
81     if format is Format.YAML:
82         return from_yaml(Topology, topology)
83     elif format is Format.JSON:
84         return from_json(Topology, topology)
85     else:
86         raise ValueError(f"Unknown format {format}")
87

```

Figure 3.6: *parseTopology* function

Above the function, different classes have been defined and tagged with the `@serde` decorator. They represent the attributes which can be defined in the network topology file. Then, each class has its own field attributes which can be optional or not. The topology object has been declared and it uses all classes defined in the script. It represents the entire network topology input file.

```

67 @serde(tagging=InternalTagging("type"))
68 class Topology:
69     network: Network
70     controllers: List[Union[ControllerLocal, ControllerRemote]]
71     switches: List[Switch]
72     hosts: List[Host]

```

Figure 3.7: Topology object

So, it means that in the input file there are:

- A Network attribute: the *autoArp* and *autoMac* flag for the network can be set.

```

9 @serde
10 class Network:
11     autoMac: bool = field(default=False)
12     autoArp: bool = field(default=False)

```

Figure 3.8: Network object

- A list of Controllers attribute: more than one controller can be defined for the same Mininet network. A controller can be local or remote. The local one has only the name attribute. Instead, the remote has name, IP address and port attributes.

```

15 @serde
16 class ControllerLocal:
17     name: str
18
19
20 @serde
21 class ControllerRemote:
22     name: str
23     ip: ipaddress.IPv4Address
24     port: int

```

Figure 3.9: Controller objects

- A list of Switches attribute: more than one switch can be emulated. A name and list of network links can be defined for each switch.

```

36 @serde
37 class Switch:
38     name: str
39     links: List[NetworkLink] = field(default_factory=list)

```

Figure 3.10: Switch object

A network link is defined by some optional parameters like bandwidth, delay, fromInterface, toInterface. The only mandatory field is the node to which the link is attached.

```

27 @serde
28 class NetworkLink:
29     node: str
30     bandwidth: Optional[int]
31     delay: Optional[str]
32     fromInterface: Optional[str]
33     toInterface: Optional[str]

```

Figure 3.11: Network link object

- A list of Hosts: an host represents an emulated node in the network. It must have a name and IP address. Optionally, a MAC address and Docker image might be defined for it. In turn, the host can deploy a list of Docker containers inside itself. A host might also have a list of network interfaces. Comnetsemu assigns one network interface by default, but the user may want to add more.

```

57 @serde
58 class Host:
59     name: str
60     ip: ipaddress.IPv4Interface
61     mac: Optional[str]
62     image: Optional[str]
63     containers: List[Container] = field(default_factory=list)
64     interfaces: List[NetworkInterface] = field(default_factory=list)

```

Figure 3.12: Host object

A container object must be defined by a name and Docker image. It may have a command to be executed when the container starts.

```

42 @serde
43 class Container:
44     name: str
45     image: str
46     cmd: Optional[str]
47     wait: Optional[bool]

```

Figure 3.13: Container object

A network interface object must have a name and IP address. Furthermore, a MAC address might be defined.

```

50 @serde
51 class NetworkInterface:
52     name: str
53     ip: ipaddress.IPv4Interface
54     mac: Optional[str]

```

Figure 3.14: Network interface object

3.3 src/topology/topologyBuilder.py

The *buildTopology* function is defined in the script. It takes the topology object and builds every component of it. Firstly, it initializes the containernet object. Then, it builds controllers, switches, hosts, network links and network interfaces. The virtual network function manager object is initialized. In the Comnetsemu environment, the VNFManager has the capability to start or stop any host. At the end, Docker containers inside hosts are built up.

```

112 def buildTopology(topology: Topology) -> Tuple[Containernet, VNFManager]:
113     logger.info("=== NETWORK ===")
114     logger.info(f"Network: {to_json(topology.network)}")
115     network = Containernet(
116         switch=node.OVSKernelSwitch,
117         autoSetMacs=topology.network.autoMac,
118         autoStaticArp=topology.network.autoArp,
119         link=mnlink.TCLink,
120     )
121
122     logger.info("=== CONTROLLERS ===")
123     buildControllers(network, topology)
124     logger.info("=== SWITCHES ===")
125     buildSwitches(network, topology)
126     logger.info("=== HOSTS ===")
127     hostInstances = buildHosts(network, topology)
128
129     logger.info("=== NETWORK LINKS ===")
130     buildNetworkLinks(network, topology)
131     logger.info("=== NETWORK INTERFACES ===")
132     buildNetworkInterfaces(hostInstances)
133
134     # Manager
135     manager = VNFManager(network)
136
137     logger.info("=== CONTAINERS ===")
138     buildContainers(manager, topology)
139
140     return (network, manager)

```

Figure 3.15: *buildTopology* function

So, the functions to build the entire network are:

- *buildControllers*: it builds all the controllers. For each controller, it checks if it is remote or local. Then, it uses the Mininet function *addController* to add it to the network.

```

10 def controllerType(controller: Union[ControllerLocal, ControllerRemote]):
11     if isinstance(controller, ControllerLocal):
12         return node.Controller
13     elif isinstance(controller, ControllerRemote):
14         return node.RemoteController
15     else:
16         raise Exception(f"Unknown class instance {type(controller)}")
17
18
19 def buildControllers(network: Containernet, topology: Topology) -> None:
20     for controller in topology.controllers:
21         logger.info(f"Controller {controller.name}: {to_json(controller)}")
22         params = {}
23
24         if isinstance(controller, ControllerRemote):
25             params["ip"] = str(controller.ip)
26             params["port"] = controller.port
27
28         network.addController(
29             controller.name, controller=controllerType(controller), **params
30         )

```

Figure 3.16: *buildControllers* function

- *buildSwitches*: it builds all the switches. For each switch, it is assigned the datapath id and OpenFlow protocol version 1.0. Then, it uses the Mininet function *addSwitch* to add it to the network.

```

33 def buildSwitches(network: Containernet, topology: Topology) -> None:
34     for index, switch in enumerate(topology.switches):
35         logger.info(f"Switch {switch.name}: {to_json(switch)}")
36         params = {"protocols": "OpenFlow10", "dpid": "%016x" % (index + 1)}
37
38         network.addSwitch(switch.name, **params)

```

Figure 3.17: *buildSwitches* function

- *buildHosts*: it builds all the hosts defining IP address, Docker image and Docker arguments. Then, it uses the Mininet function *addDockerHost* to add it to the network.

```

41 def buildHosts(network: Containernet, topology: Topology) -> Dict[DockeHost, Host]:
42     hostInstances: Dict[DockeHost, Host] = {}
43
44     for host in topology.hosts:
45         logger.info(f"Host {host.name}: {to_json(host)}")
46         docker_args = {"hostname": host.name, "pid_mode": "host"}
47         params = {
48             "ip": host.ip.with_prefixlen,
49             "dimage": host.image if host.image else "dev_host:latest",
50             "inNamespace": True,
51             "docker_args": docker_args,
52         }
53
54         if host.mac:
55             params["mac"] = host.mac
56
57         instance = network.addDockerHost(host.name, **params)
58
59         # Add instance
60         hostInstances[instance] = host
61
62     return hostInstances

```

Figure 3.18: *buildHosts* function

- *buildNetworkLinks*: it builds all the network links. For each switch, the links between it and the nodes are built up. Then, it uses the Mininet function *addLink* to add it to the network.

```

65 def buildNetworkLinks(network: ContainerNet, topology: Topology) -> None:
66     for switch in topology.switches:
67         logger.info(f"Switch {switch.name}:")
68
69         for link in switch.links:
70             logger.info(f" {to_json(link)}")
71             params = {}
72
73             if link.bandwidth:
74                 params["bw"] = link.bandwidth
75             if link.delay:
76                 params["delay"] = link.delay
77             if link.fromInterface and link.toInterface:
78                 params["intfName1"] = link.fromInterface
79                 params["intfName2"] = link.toInterface
80
81             network.addLink(switch.name, link.node, **params)

```

Figure 3.19: *buildNetworkLinks* function

- *buildNetworkInterfaces*: it builds any extra network interface for any host. It uses the *ip addr* Linux command to assign the IP address to the interface. Moreover, it uses the *macchanger* command to assign the MAC address to the interface.

```

84 def buildNetworkInterfaces(hostInstances: Dict[DockeHost, Host]) -> None:
85     for instance, host in hostInstances.items():
86         logger.info(f"Host {host.name}:")
87
88         for interface in host.interfaces:
89             logger.info(f" {to_json(interface)}")
90
91             instance.cmd(f"ip addr add {interface.ip} dev {interface.name}")
92             if interface.mac:
93                 instance.cmd(f"macchanger -m {interface.mac} {interface.name}")

```

Figure 3.20: *buildNetworkInterfaces* function

- *buildContainers*: it builds Docker containers inside any host. It allows to do the Docker in Docker virtualization. For each host, all the declared containers are created using the VFNManager.

```

96 def buildContainers(manager: VNFManager, topology: Topology) -> None:
97     for host in topology.hosts:
98         logger.info(f"Host {host.name}:")
99
100         for container in host.containers:
101             logger.info(f" {to_json(container)}")
102
103             manager.addContainer(
104                 name=container.name,
105                 dhost=host.name,
106                 dimage=container.image,
107                 dcmd=container.cmd if container.cmd else "",
108                 wait=container.wait if container.wait else False,
109             )

```

Figure 3.21: *buildContainers* function

At the end of the script, the *cleanTopology* function has been defined. For each host, it removes the inside Docker containers. Then, it stops the network and the VNFManager.

```

143 def cleanTopology(
144     topology: Topology, network: Containernet, manager: VNFMManager
145 ) -> None:
146     logger.info("=== CLEANING ===")
147     # Container
148     for host in topology.hosts:
149         for container in host.containers:
150             logger.info(f"Container: {container.name}")
151             manager.removeContainer(container.name, True)
152     # Stop network
153     network.stop()
154     manager.stop()

```

Figure 3.22: *cleanTopology* function

3.4 Build a network topology

The command to build the network topology defined in a YAML/JSON files is:

```
$ sudo python3 topology.py --file scenarios/1/topology.yaml
```

The output should be similar to the figure 3.23.

[illegible]

Figure 3.23: Building topology output

Chapter 4

Tested Scenarios

Some possible tested scenarios and their results are reported in this chapter. First of all, the deployed service and how the migration works are described to understand better why the migration is dynamic, stateful and transparent. After that, each scenario has been defined by:

- topology.yaml: YAML file which implements the network topology with the new topology building system demonstrated in the chapter 3.
- OpenFlow: a software component which implements the OpenFlow protocol. It is needed to define slices and flows. FlowVisor [8] and ovs-ofctl [14] have been used during testing.
- RYU Controller: a python script which implements a RYU controller. It is used for handling the traffic within slices and flows. The controller might also add, update, modify or delete flows redirecting dynamically the network traffic.

Any host in Mininet network is deployed as a Docker container. So, three custom Docker images located at *src/docker* were developed:

- Dockerfile.dev_host: it is the default Mininet Docker image with network tools.
- Dockerfile.dev_server: it deploys the developed service to be migrated.
- Dockerfile.flowvisor: it installs and deploys FlowVisor for slicing the network.

4.1 Tested Service

The service is a python Flask [7] web application located at *src/server.py* and deployed in the network through the Dockerfile.dev_server. The application increments a counter every time a user performs an HTTP post request to the API */api/counter*.

```
6  # Vars
7  counter = 0
8  enable = False
9
10 # Flask app(s)
11 app = Flask(__name__)
```

Figure 4.1: Flask counter

The web application exposes four APIs:

- `/api/counter` [GET]: it returns the current counter value.

```
26 # APP
27 @app.route("/api/counter", methods=["GET"])
28 def get_counter():
29     """
30     Return counter.
31     """
32
33     global counter
34     return jsonify(counter=counter)
```

Figure 4.2: Flask `/api/counter` [GET]

- `/api/counter` [POST]: it increments and returns the counter value.

```
37 @app.route("/api/counter", methods=["POST"])
38 def post_counter():
39     """
40     Increment and return counter.
41     """
42
43     global counter
44     counter += 1
45     return jsonify(counter=counter)
```

Figure 4.3: Flask `/api/counter` [POST]

- `/api/admin/disable` [POST]: it logically disables the service setting to *False* the global variable *enable* and returns the current counter value.

```
71 @app.route("/api/admin/disable", methods=["POST"])
72 def disable():
73     """
74     Disable service.
75     """
76
77     global enable
78     enable = False
79     return jsonify(counter=counter)
```

Figure 4.4: Flask `/api/admin/disable` [POST]

- `/api/admin/migrate` [POST]: it performs the migration detailed in the section 4.2. It retrieves the IP address of the current enabled server from the body parameter. Then, it disables it and gets the current counter value calling the API `/api/admin/disable`. At the end, it stores the current counter value and set the global variable *enable* to *True* to become the new enabled server.


```

48 # APPMGR
49 @app.route("/api/admin/migrate", methods=["POST"])
50 def migrate():
51     """
52     Migrate service and return counter.
53     """
54
55     global counter, enable
56
57     # Request body
58     body = request.get_json(force=True)
59     # Server
60     server = body["server"]
61     # Disable other server & obtain counter
62     response = requests.post(f"{server}/api/admin/disable").json()
63     # Set counter value
64     counter = response["counter"]
65     # Enable this server
66     enable = True
67
68     return jsonify(counter=counter)

```

Figure 4.5: Flask `/api/admin/migrate` [POST]

4.2 How Migration Works

Firstly, there are at least two running Flask web services in every scenario and only one is enabled to receive HTTP requests. Moreover, if a user performs a GET or POST request on the endpoint `/api/counter`, it gets a 503 HTTP status code, because it is intercepted and blocked by the middleware reported in figure 4.6.

```

14 @app.before_request
15 def before_request_func():
16     global enable
17
18     # Check URL admin request
19     admin = re.search(r"\badmin\b", request.url)
20
21     # Return 503 if disabled
22     if not enable and not admin:
23         return "Service Unavailable", 503

```

Figure 4.6: Flask middleware

Ideally, the migration should be performed by a network administrator and it works as following:

1. The administrator performs an HTTP post request at `/api/admin/migrate` to the disabled server. He has to provide the IP address of the enabled server in the body following the JSON format.
2. The disabled server retrieves the IP address and performs an HTTP post request at `[Enabled Server IP Address]/api/admin/disable`.
3. The enabled server sets its own global variable `enable` to false and returns the counter current value in the response.
4. The disabled server stores the counter value and sets to true its own global variable `enable`.

4.3 Scenario 1

The Scenario 1 is located in the folder `src/scenarios/1` and the network topology is reported in figure 4.7.

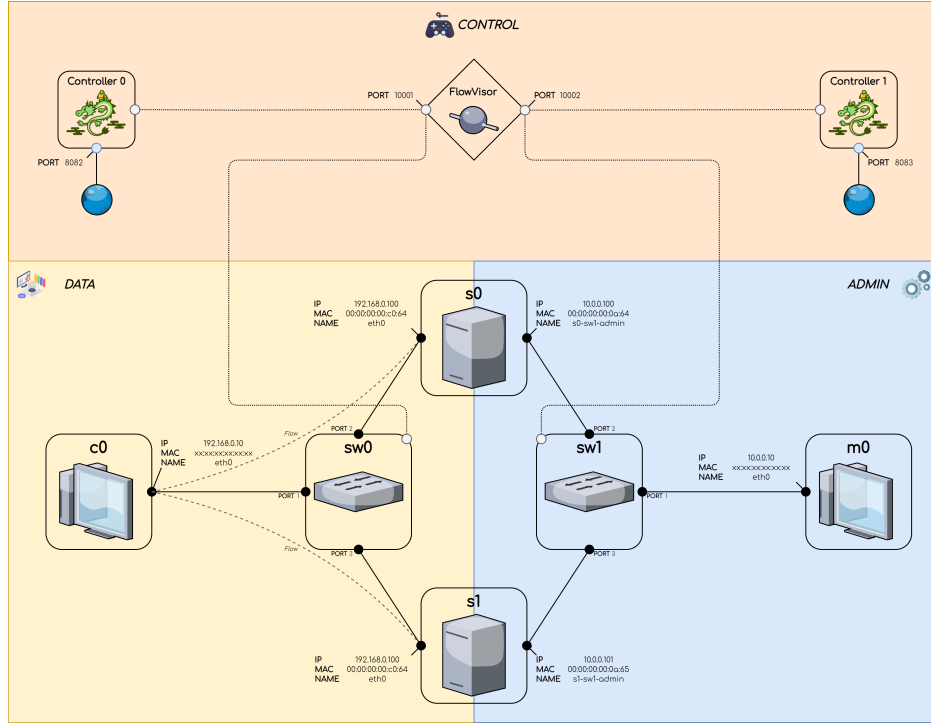


Figure 4.7: Network topology

There are several devices in the network:

- Client *c0*: it performs the HTTP requests to the enabled server.
- Servers *s0* and *s1*: they implement the Flask web application. By default, *s0* is the enabled server and *s1* is disabled. After the migration, *s1* will become enabled getting the current counter value from *s0*.
- Switch *sw0*: it connects *c0* with the two servers *s0* and *s1*.
- Manager *m0*: it emulates the network administrator and executes the command for the migration.
- Switch *sw1*: it connects *m0* with the two servers *s0* and *s1*.
- Controller 0: it handles the network traffic of the *data* slice telling to the switch *sw0* how to forward the traffic.
- Controller 1: it handles the network traffic of the *admin* slice telling to the switch *sw1* how to forward the traffic.

So, the OpenFlow protocol defines slices and flows and it is implemented by FlowVisor in the scenario. The *data* slice were created to emulate the traffic generated by a standard user. Instead, the *admin* slice emulates the administrator traffic to perform the service migration. Slices and flows are defined by the script *src/scenarios/1/flowvisor.sh*.

```

47 # FlowVisor admin slice & flow
48 INFO "Creating FlowVisor admin slice"
49 fvctl_exec add-slice --password=password slice_service_migration_admin tcp:localhost:10002 admin@slice_service_migration_admin
50 INFO "Creating FlowVisor admin flow"
51 fvctl_exec add-flowspace admin 2 1 any slice_service_migration_admin=7
52
53 # FlowVisor data slice
54 INFO "Creating FlowVisor data slice"
55 fvctl_exec add-slice --password=password slice_service_migration tcp:localhost:10001 admin@slice_service_migration
56
57 # FlowVisor data flow
58 INFO "Creating FlowVisor data flow"
59 fvctl_exec add-flowspace dpid1-c0 1 1 in_port=1 slice_service_migration=7
60 fvctl_exec add-flowspace dpid1-s 1 1 in_port="{SERVERS_PORT[IDX_SERVER]}" slice_service_migration=7

```

Figure 4.8: Slices and flows

The approach for the migration is to adjust dynamically the flow between the switch *sw0* and servers in order to always connect the client to the enabled server. As reported in figure 4.8, one flow was created for the *admin* slice and two for the *data* slice. So, in the *data* slice there is one flow to connect the client *c0* to the switch *sw0* and the other one is to connect the switch *sw0* to the enabled server. The latter is a dynamic flow, because the OpenFlow attribute *in_port* is automatically changed towards the enabled server for performing the service migration. Hence, the migration is straightforward in theory, but the implementation is a bit tricky. Firstly, FlowVisor needs a container itself for running and the script *src/scripts/flowvisor.sh* deploys a Docker container for FlowVisor. Secondly, the manager *m0* needs to send the command to the disabled server for starting the migration. So, the script *src/scripts/flowvisor.py* is a Flask web server which accepts HTTP post request at */api/migrate* and sends the command to *m0* using the Docker command line.

```

11 @app.route("/api/migrate", methods=['POST'])
12 def migrate():
13     """
14     Migrate service via Docker.
15     """
16
17     # Request body
18     body = request.get_json(force=True)
19     # From server
20     from_server = body["from"]
21     # To server
22     to_server = body["to"]
23
24     # Call Docker
25     subprocess.call(
26         f'docker exec m0 curl -X POST -H "Content-Type:application/json" -d \'{{"server": "http://{from_server}" }}\' {to_server}/api/admin/migrate',
27         shell=True,
28     )
29
30     return ("", 200)

```

Figure 4.9: Flask *flowvisor.py*

The web server is run by the *src/scripts/flowvisor.sh* before the FlowVisor container deployment.

```

61 # Start server
62 INFO "Starting Docker server..."
63 python3 "${_DIRNAME}/flowvisor.py" --port $PORT > /dev/null 2>&1 &
64 SERVER_PID=$(jobs -p)
65 readonly SERVER_PID
66 INFO "Docker server started with PID $SERVER_PID"

```

Figure 4.10: Run Flask *flowvisor.py*

Thirdly, the Controller 0 has to keep tracking of the new flow. The Controller 0 is a simple self-learning controller and for each datapath id, the controller maps the MAC address to the attribute *in_port*. It is able to automatically learn the mapping between MAC address and the OpenFlow attribute *in_port* using the OpenFlow Flood.

```

94         # Learn mac address
95         self.mac_to_port.setdefault(dpid, {})
96         self.mac_to_port[dpid][src] = in_port
97
98         # Find out_port
99         if dst in self.mac_to_port[dpid]:
100             out_port = self.mac_to_port[dpid][dst]
101         else:
102             out_port = ofproto.OFPP_FLOOD
103
104         actions = [parser.OFPAActionOutput(out_port)]

```

Figure 4.11: Self-learning controller

As soon as the controller receives packet, it tries to lookup the *in_port* of the destination MAC address specified in the packet. If it is able to find the *in_port*, it sets the OpenFlow actions attribute *out_port* to it. Otherwise, it starts an OpenFlow Flood to discover the unknown *in_port* flooding the network. The controller self-learning code is reported in figure 4.11. When the flow changes, the mapping stored by the controller has to be updated as well. The script *src/scenarios/1/migrate.py* is a Flask web server run by the controller in a new python thread to handle incoming HTTP request for updating the mapping. As soon as the controller is launched, it spawns a new thread for the web server as reported in figure 4.12.

```

15 class Controller(app_manager.RyuApp):
16     OFP_VERSIONS = [ofproto_v1_0.OFP_VERSION]
17
18     def __init__(self, *args, **kwargs):
19         super(Controller, self).__init__(*args, **kwargs)
20         CONF = cfg.CONF
21         CONF.register_opts(
22             [
23                 cfg.IntOpt("port", default=None),
24             ]
25         )
26
27         self.mac_to_port = {}
28
29         if CONF.port:
30             migrator_port = int(CONF.port)
31             self.thread_migration = threading.Thread(
32                 target=self.thread_migration_cb,
33                 daemon=True,
34                 kwargs={"port": migrator_port},
35             )
36             self.thread_migration.start()
37
38     def thread_migration_cb(self, port: int):
39         migrator.start(port, self.migration_cb)
40
41     def migration_cb(self, dpid: int, mac: str, port: int):
42         self.mac_to_port[dpid][mac] = port

```

Figure 4.12: Spawn *migrator.py*

Hence, the *migrator.py* accepts HTTP post request at */api/migrate* and call the method *migration_cb* to update the mapping. The datapath id, MAC address and the new *in_port* has to be passed in the body of the request.

```

9  @app.route("/api/migrate", methods=["POST"])
10 def migrate():
11     """
12     Migrate service.
13     """
14
15     # Request body
16     body = request.get_json(force=True)
17     # Switch dpid
18     dpid = int(body["dpid"])
19     # Server mac
20     mac = body["mac"]
21     # Server port
22     port = int(body["port"])
23
24     # Notify callback
25     if cb:
26         cb(dpid, mac, port)
27
28     return ("", 200)
29
30
31 def start(port: int, callback):
32     global cb
33     cb = callback
34     app.run(port=port)

```

Figure 4.13: *migrator.py* API

To sum up, there is a while loop for migration in the script *src/scenarios/1/flowvisor.sh*. If the user press Enter:

- An HTTP post request has been performed to the *src/script/flowvisor.py* Flask web server to launch the migration command from the manager *m0*.
- An HTTP post request has been performed to the *src/scenarios/1/migrator.py* Flask web server to update the controller mapping.

```

62 # Migration loop
63 while read -n1 -r -p "Press 'Enter' to migrate or 'q' to exit" && [[ $REPLY != q ]]; do
64     # Old ip
65     OLD_IP=${SERVERS_IP[IDX_SERVER]}
66     # Old port
67     OLD_PORT=${SERVERS_PORT[IDX_SERVER]}
68     # Update idx server
69     IDX_SERVER=$((next_idx_server "$IDX_SERVER"))
70     # New ip
71     NEW_IP=${SERVERS_IP[IDX_SERVER]}
72     # New port
73     NEW_PORT=${SERVERS_PORT[IDX_SERVER]}
74
75     INFO "Migrating from { ip: $OLD_IP, port: $OLD_PORT } to { ip: $NEW_IP, port: $NEW_PORT }"
76
77     # Docker manager
78     curl -X POST -H "Content-Type:application/json" -d "{ \"from\": \"$OLD_IP\", \"to\": \"$NEW_IP\" }" localhost:12345/api/migrate
79
80     # Controller flow
81     curl -X POST -H "Content-Type:application/json" -d "{ \"dpid\": \"$1\", \"mac\": \"$SERVER_MAC\", \"port\": \"$NEW_PORT\" }" localhost:9876/api/migrate
82
83     # FlowVisor
84     fvctl_exec update-flowspace --match-in_port="$NEW_PORT" dpid1-s
85     fvctl_stop
86     fvctl_start
87
88     INFO "Successfully migrated from { ip: $OLD_IP, port: $OLD_PORT } to { ip: $NEW_IP, port: $NEW_PORT }"
89 done

```

Figure 4.14: While loop migration

4.3.1 Running the Scenario

Repeat the following commands three times, because three terminals in the Comnetsemu machine are needed.

```

$ cd SVMN/comnetsemu
$ vagrant ssh
$ cd comnetsemu/app/morphing_slices

```

Then, follow the steps:

1. Generate the network from the topology.yaml file in the first terminal:

```
$ sudo python3 topology.py --file scenarios/1/topology.yaml
```

The output should be similar to figure 4.15 executing the Mininet command *dump*.

```
mininet> dump
<DockerHost m0: m0-eth0:10.0.0.10 pid=2244>
<DockerHost c0: c0-eth0:192.168.0.10 pid=2361>
<DockerHost s0: s0-eth0:192.168.0.100,s0-sw1-admin:None pid=2462>
<DockerHost s1: s1-eth0:192.168.0.100,s1-sw1-admin:None pid=2568>
<OVSSwitch sw0: lo:127.0.0.1,sw0-eth1:None,sw0-eth2:None,sw0-eth3:None pid=2184>
<OVSSwitch sw1: lo:127.0.0.1,sw1-eth1:None,sw1-s0-admin:None,sw1-s1-admin:None pid=2187>
<RemoteController controller0: 127.0.0.1:6633 pid=2177>
```

Figure 4.15: Mininet dump

2. In the second terminal, start the FlowVisor container attaching the folder *scenarios/1* as volume:

```
$ scripts/flowvisor.sh --volume scenarios/1
```

```
vagrant@comnetsemu:~/comnetsemu/app/morphing_slices$ scripts/flowvisor.sh --volume scenarios/1
[2022-06-26 15:18:26.735][INFO ][main:62 ] Starting Docker server...
[2022-06-26 15:18:26.742][INFO ][main:66 ] Docker server started with PID 10045
[2022-06-26 15:18:26.747][INFO ][main:69 ] Run Flowvisor image 'flowvisor:latest'
[2022-06-26 15:18:26.754][INFO ][main:71 ] Mounted shared volume /home/vagrant/comnetsemu/app/morphing_slices/scenarios/1
root@comnetsemu:~/flowvisor#
```

Figure 4.16: Run FlowVisor container

The script for creating the Docker container is located in *src/scripts/flowvisor.sh*.

Then, run the script *src/scenarios/1/flowvisor.sh* in the FlowVisor container:

```
$ ./flowvisor.sh
```

```
root@comnetsemu:~/flowvisor# ./flowvisor.sh
[2022-06-26 15:19:01.367][INFO ][fvctl_start:31 ] Starting FlowVisor service...
Starting flowvisor with the configuration stored in DB
If DB unpopulated, load config using 'fvconfig load config.json'
[2022-06-26 15:19:04.390][INFO ][fvctl_start:34 ] FlowVisor service started
[2022-06-26 15:19:04.396][INFO ][main:48 ] Creating FlowVisor admin slice
Slice slice_service_migration_admin was successfully created
[2022-06-26 15:19:05.116][INFO ][main:50 ] Creating FlowVisor admin flow
[2022-06-26 15:19:05.206][INFO ][main:54 ] Creating FlowVisor data slice
Slice slice_service_migration was successfully created
[2022-06-26 15:19:05.458][INFO ][main:58 ] Creating FlowVisor data flow
FlowSpace dpid1-c0 was added with request id 1.
FlowSpace dpid1-s was added with request id 2.
Press 'Enter' to migrate or 'q' to exit
```

Figure 4.17: Run *flowvisor.sh* migration script

The script is already in the container, because the folder has been mounted in the previous step.

3. In the third terminal, run the two RYU controllers:

```
$ parallel -j 2 --ungroup ::: \
'scripts/ryu.sh --controller scenarios/1/controller.py --ofport 10001 \
--port 8082 --config scenarios/1/controller.cfg' \
'scripts/ryu.sh --controller scenarios/1/controller.py --ofport 10002 \
--port 8083'
```

```
root@ubuntu:~/controller# ./scripts/ryu.sh --controller scenarios/1/controller.py --ofport 10001 --port 8082 --config scenarios/1/controller.cfg 'sc
[2022-06-26 15:21:09.789] INFO [main:76] Running ryu: ( controller: scenarios/1/controller.py, ofport: 10001, port: 8082, config: )
[2022-06-26 15:21:09.792] INFO [main:76] Running ryu: ( controller: scenarios/1/controller.py, ofport: 10001, port: 8082, config: scenarios/1/controller.cfg )
Loading app /usr/local/lib/python3.8/dist-packages/ryu/app/gui_topology/gui_topology.py
Loading app /usr/local/lib/python3.8/dist-packages/ryu/app/gui_topology/gui_topology.py
Loading app scenarios/1/controller.py
Loading app scenarios/1/controller.py
Loading app ryu.app.ws_topology
Loading app ryu.app.ws_topology
Loading app ryu.app.rest_topology
Loading app ryu.app.rest_topology
Loading app ryu.controller.ofp_handler
Loading app ryu.controller.ofp_handler
Creating context wsgi
Instantiating app None of Switches
Instantiating app None of DPSet
Creating context dpidet
Instantiating app /usr/local/lib/python3.8/dist-packages/ryu/app/gui_topology/gui_topology.py of GUIserverApp
Instantiating app scenarios/1/controller.py of Controller
Instantiating app ryu.app.ws_topology of WebsocketTopology
Instantiating app ryu.app.rest_topology of TopologyAPI
Instantiating app ryu.app.ofctl_rest of RestStatsAPI
Instantiating app ryu.controller.ofp_handler of OFPHandler
[10292] wsgi starting up on http://0.0.0.0:8082
Loading app ryu.app.rest_topology
Loading app ryu.app.ws_topology
Loading app ryu.controller.ofp_handler
Loading app ryu.controller.ofp_handler
Creating context wsgi
Instantiating app None of DPSet
Creating context dpidet
Instantiating app None of Switches
Creating context wsgi
Instantiating app /usr/local/lib/python3.8/dist-packages/ryu/app/gui_topology/gui_topology.py of GUIserverApp
Instantiating app scenarios/1/controller.py of Controller
+ Sending flask app 'relarguer' (lazy loading)
+ Environment: production
+ WARNING: This is a development server. Do not use it in a production deployment.
+ Use a production WSGI server instead.
+ Debug mode: off
+ Running on http://127.0.0.1:8082 (Press CTRL+C to quit)
Instantiating app ryu.app.ofctl_rest of RestStatsAPI
Instantiating app ryu.app.ws_topology of WebsocketTopology
Instantiating app ryu.app.rest_topology of TopologyAPI
Instantiating app ryu.controller.ofp_handler of OFPHandler
[10292] wsgi starting up on http://0.0.0.0:8082
OFPMatchFeatures: { 'cookie': 2, 'buffers': 0, 'tables': 256, 'capabilities': 199, 'ports': [ { port: 80334, name: sw1, hw_addr: be:23:7a:75:e9:08 }, { port: 1, name: sw1-eth1, hw_addr: 3e:7a:4a:a8:02:0f }, { port: 2, na
sw1-sw-eth1, hw_addr: 0e:ad:39:dd:06 }, { port: 3, name: sw1-i-sw-hw, hw_addr: 76:97:1d:1b:3f:a7 } ] }
OFPMatchFeatures: { dpid: 1, buffers: 0, tables: 256, capabilities: 199, 'ports': [ { port: 1, name: sw0-eth1, hw_addr: 56:00:03:d7:0b:98 }, { port: 2, name: sw0-eth2, hw_addr: f6:a5:3b:02:17:4e } ] }
watches: { dpid: 1 }
watcher: { dpid: 1 }
```

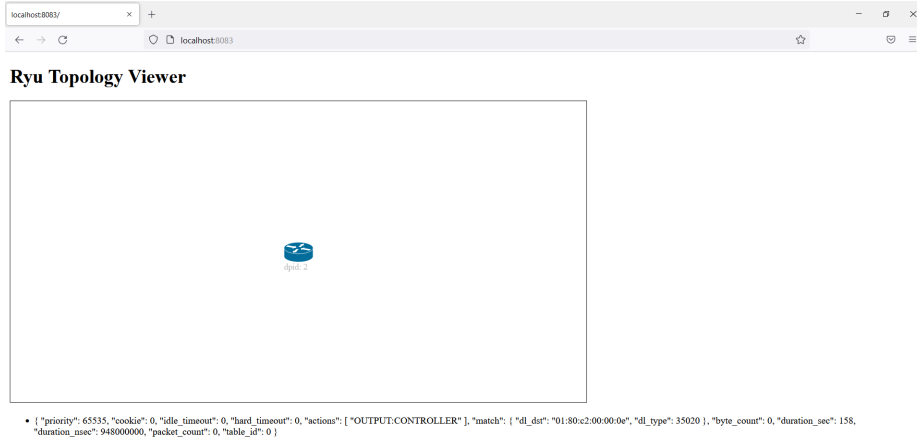
Figure 4.18: Run controllers

The RYU Graphical User Interfaces can be accessed from the browser at the link *http://localhost:8082* and *http://localhost:8083*.

The RYU GUI shows the datapath handled by the controller. So, the Controller 0 handles the datapath 1 related to the switch *sw0* as reported in the figure 4.19.

Figure 4.19: RYU GUI at *http://localhost:8082*

The Controller 1 handles the datapath 2 related to the switch *sw1* as reported in the figure 4.20.

Figure 4.20: RYU GUI at <http://localhost:8083>

Any request can be performed to the enabled server in the terminal where the network topology has been deployed. So, the client can increment the counter.

```
$ c0 curl -X POST 192.168.0.100/api/counter
```

```
mininet> c0 curl -X POST 192.168.0.100/api/counter
{
  "counter": 1
}
mininet> c0 curl -X POST 192.168.0.100/api/counter
{
  "counter": 2
}
mininet> c0 curl -X POST 192.168.0.100/api/counter
{
  "counter": 3
}
mininet> c0 curl -X POST 192.168.0.100/api/counter
{
  "counter": 4
}
```

Figure 4.21: Client increments counter

At the beginning, the server *s0* is enabled and its *in_port* is the 2. Indeed, *s0* is replying to the client checking from the controller logs reported in the figure 4.22. Instead, the server *s1* is disabled and its *in_port* is the 3.

```
OFPPacketIn: { dpid: 1, src: 00:00:00:00:c0:64, dst: d6:08:d8:eb:1c:df, in_port: 2 }
OFPPacketIn: { dpid: 1, src: d6:08:d8:eb:1c:df, dst: 00:00:00:00:c0:64, in_port: 1 }
OFPPacketIn: { dpid: 1, src: 00:00:00:00:c0:64, dst: d6:08:d8:eb:1c:df, in_port: 2 }
OFPPacketIn: { dpid: 1, src: d6:08:d8:eb:1c:df, dst: 00:00:00:00:c0:64, in_port: 1 }
OFPPacketIn: { dpid: 1, src: d6:08:d8:eb:1c:df, dst: 00:00:00:00:c0:64, in_port: 1 }
OFPPacketIn: { dpid: 1, src: 00:00:00:00:c0:64, dst: d6:08:d8:eb:1c:df, in_port: 2 }
OFPPacketIn: { dpid: 1, src: 00:00:00:00:c0:64, dst: d6:08:d8:eb:1c:df, in_port: 2 }
OFPPacketIn: { dpid: 1, src: 00:00:00:00:c0:64, dst: d6:08:d8:eb:1c:df, in_port: 2 }
OFPPacketIn: { dpid: 1, src: d6:08:d8:eb:1c:df, dst: 00:00:00:00:c0:64, in_port: 1 }
OFPPacketIn: { dpid: 1, src: d6:08:d8:eb:1c:df, dst: 00:00:00:00:c0:64, in_port: 1 }
OFPPacketIn: { dpid: 1, src: 00:00:00:00:c0:64, dst: d6:08:d8:eb:1c:df, in_port: 2 }
OFPPacketIn: { dpid: 1, src: 00:00:00:00:c0:64, dst: d6:08:d8:eb:1c:df, in_port: 2 }
OFPPacketIn: { dpid: 1, src: d6:08:d8:eb:1c:df, dst: 00:00:00:00:c0:64, in_port: 1 }
OFPPacketIn: { dpid: 1, src: d6:08:d8:eb:1c:df, dst: 00:00:00:00:c0:64, in_port: 1 }
OFPPacketIn: { dpid: 1, src: 00:00:00:00:c0:64, dst: d6:08:d8:eb:1c:df, in_port: 2 }
OFPPacketIn: { dpid: 1, src: 00:00:00:00:c0:64, dst: d6:08:d8:eb:1c:df, in_port: 2 }
OFPPacketIn: { dpid: 1, src: d6:08:d8:eb:1c:df, dst: 00:00:00:00:c0:64, in_port: 1 }
OFPPacketIn: { dpid: 1, src: d6:08:d8:eb:1c:df, dst: 00:00:00:00:c0:64, in_port: 1 }
```

Figure 4.22: *s0* replies to client

The migration can be done just pressing *Enter* in the FlowVisor terminal.

```
root@comnetsemu:~/flowvisor# ./flowvisor.sh
[2022-06-26 15:19:01.367][INFO ][fvctl_start:31 ] Starting FlowVisor service...
Starting flowvisor with the configuration stored in DB
If DB unpopulated, load config using 'fvconfig load config.json'
[2022-06-26 15:19:04.390][INFO ][fvctl_start:34 ] FlowVisor service started
[2022-06-26 15:19:04.396][INFO ][main:48 ] Creating FlowVisor admin slice
Slice slice_service_migration_admin was successfully created
[2022-06-26 15:19:05.116][INFO ][main:50 ] Creating FlowVisor admin flow
[2022-06-26 15:19:05.206][INFO ][main:54 ] Creating FlowVisor data slice
Slice slice_service_migration was successfully created
[2022-06-26 15:19:05.458][INFO ][main:58 ] Creating FlowVisor data flow
FlowSpace dpid1-c0 was added with request id 1.
FlowSpace dpid1-s was added with request id 2.
Press 'Enter' to migrate or 'q' to exit
```

Figure 4.23: FlowVisor migration

The client can continue to increment the counter where it has left without noticing the migration.

```
mininet> c0 curl -X POST 192.168.0.100/api/counter
{
  "counter": 1
}
mininet> c0 curl -X POST 192.168.0.100/api/counter
{
  "counter": 2
}
mininet> c0 curl -X POST 192.168.0.100/api/counter
{
  "counter": 3
}
mininet> c0 curl -X POST 192.168.0.100/api/counter
{
  "counter": 4
}
mininet> c0 curl -X POST 192.168.0.100/api/counter
{
  "counter": 5
}
mininet> c0 curl -X POST 192.168.0.100/api/counter
{
  "counter": 6
}
mininet> c0 curl -X POST 192.168.0.100/api/counter
{
  "counter": 7
}
```

Figure 4.24: *c0* increments the counter after migration

However, the server *s1* is the enabled and the *s0* is disabled. It can be checked from the controller logs as reported in figure 4.25.

```
OFPPacketIn: { dpid: 1, src: d6:08:d8:eb:1c:df, dst: 00:00:00:00:c0:64, in_port: 1 }
OFPPacketIn: { dpid: 1, src: 00:00:00:00:c0:64, dst: ff:ff:ff:ff:ff:ff, in_port: 3 }
OFPPacketIn: { dpid: 1, src: d6:08:d8:eb:1c:df, dst: 00:00:00:00:c0:64, in_port: 1 }
OFPPacketIn: { dpid: 1, src: 00:00:00:00:c0:64, dst: d6:08:d8:eb:1c:df, in_port: 3 }
OFPPacketIn: { dpid: 1, src: d6:08:d8:eb:1c:df, dst: 00:00:00:00:c0:64, in_port: 1 }
OFPPacketIn: { dpid: 1, src: d6:08:d8:eb:1c:df, dst: 00:00:00:00:c0:64, in_port: 1 }
OFPPacketIn: { dpid: 1, src: 00:00:00:00:c0:64, dst: d6:08:d8:eb:1c:df, in_port: 3 }
OFPPacketIn: { dpid: 1, src: 00:00:00:00:c0:64, dst: d6:08:d8:eb:1c:df, in_port: 3 }
OFPPacketIn: { dpid: 1, src: 00:00:00:00:c0:64, dst: d6:08:d8:eb:1c:df, in_port: 3 }
OFPPacketIn: { dpid: 1, src: d6:08:d8:eb:1c:df, dst: 00:00:00:00:c0:64, in_port: 1 }
OFPPacketIn: { dpid: 1, src: 00:00:00:00:c0:64, dst: d6:08:d8:eb:1c:df, in_port: 3 }
OFPPacketIn: { dpid: 1, src: 00:00:00:00:c0:64, dst: d6:08:d8:eb:1c:df, in_port: 3 }
OFPPacketIn: { dpid: 1, src: 00:00:00:00:c0:64, dst: d6:08:d8:eb:1c:df, in_port: 3 }
OFPPacketIn: { dpid: 1, src: d6:08:d8:eb:1c:df, dst: 00:00:00:00:c0:64, in_port: 1 }
OFPPacketIn: { dpid: 1, src: 00:00:00:00:c0:64, dst: d6:08:d8:eb:1c:df, in_port: 3 }
OFPPacketIn: { dpid: 1, src: d6:08:d8:eb:1c:df, dst: 00:00:00:00:c0:64, in_port: 1 }
OFPPacketIn: { dpid: 1, src: 00:00:00:00:c0:64, dst: d6:08:d8:eb:1c:df, in_port: 3 }
OFPPacketIn: { dpid: 1, src: 00:00:00:00:c0:64, dst: d6:08:d8:eb:1c:df, in_port: 3 }
OFPPacketIn: { dpid: 1, src: d6:08:d8:eb:1c:df, dst: 00:00:00:00:c0:64, in_port: 1 }
OFPPacketIn: { dpid: 1, src: d6:08:d8:eb:1c:df, dst: 00:00:00:00:c0:64, in_port: 1 }
OFPPacketIn: { dpid: 1, src: 00:00:00:00:c0:64, dst: d6:08:d8:eb:1c:df, in_port: 3 }
```

Figure 4.25: *s1* replies to client

4.4 Scenario 2

The Scenario 2 is located in the folder `src/scenarios/2` and the network topology is reported in figure 4.26.

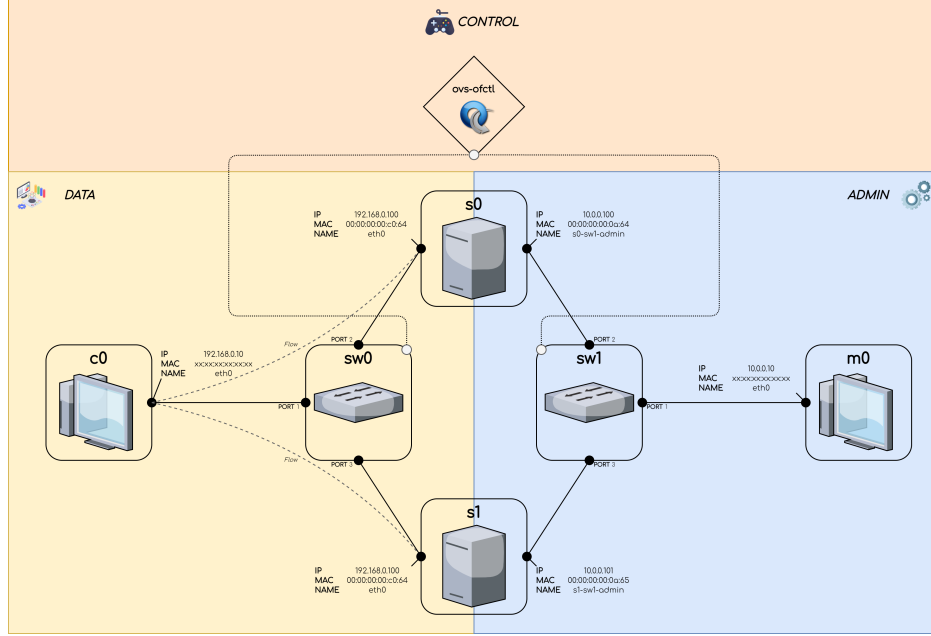


Figure 4.26: Network topology

There are several devices in the network:

- Client *c0*: it performs the HTTP requests to the enabled server.
- Servers *s0* and *s1*: they implement the Flask web application. By default, *s0* is the enabled server and *s1* is disabled. After the migration, *s1* will become enabled getting the current counter value from *s0*.
- Switch *sw0*: it connects *c0* with the two servers *s0* and *s1*.
- Manager *m0*: it emulates the network administrator and executes the command for the migration.
- Switch *sw1*: it connects *m0* with the two servers *s0* and *s1*.

The OpenFlow protocol is implemented by `ovs-ofctl` in the scenario which defines flows for both *data* and *admin* subnets. Flows are defined by the script `src/scenarios/2/openflow.sh`.

```

26 # OpenFlow admin flow
27 INFO "Creating OpenFlow admin flow"
28 ofctl_exec add-flow sw1 actions=normal
29
30 # OpenFlow data flow
31 INFO "Creating OpenFlow data flow"
32 ofctl_exec add-flow sw0 in_port=1,actions=output:"${SERVERS_PORT[IDX_SERVER]}"
33 ofctl_exec add-flow sw0 in_port=2,actions=output:1
34 ofctl_exec add-flow sw0 in_port=3,actions=output:1

```

Figure 4.27: ovs-ofctl flows

As shown in figure 4.27, one flow is created for the *admin* and three flows for the *data* defined by the OpenFlow attribute *in_port* and the action *out_port*. The migration works adjusting dynamically the flow between the switch *sw0* and servers. So, there is a while loop to wait the user input to perform the migration. After pressing *Enter*, the *out_port* of the flow will be changed and the manager *m0* will perform the migration between servers *s0* and *s1* using Docker command line.

```

36 # Migration loop
37 while read -n1 -r -p "Press 'Enter' to migrate or 'q' to exit" && [[ $REPLY != q ]]; do
38     # Old ip
39     OLD_IP=${SERVERS_IP[IDX_SERVER]}
40     # Old port
41     OLD_PORT=${SERVERS_PORT[IDX_SERVER]}
42     # Update idx server
43     IDX_SERVER=$((next_idx_server $IDX_SERVER))
44     # New ip
45     NEW_IP=${SERVERS_IP[IDX_SERVER]}
46     # New port
47     NEW_PORT=${SERVERS_PORT[IDX_SERVER]}
48
49     INFO "Migrating from { ip: $OLD_IP, port: $OLD_PORT } to { ip: $NEW_IP, port: $NEW_PORT }"
50
51     # Docker manager
52     docker exec --detach m0 curl -X POST -H "Content-Type:application/json" -d "{ \"server\": \"http://$OLD_IP\" }" "$NEW_IP/api/admin/migrate"
53
54     # OpenFlow data flow
55     ofctl_exec mod-flows sw0 in_port=1,actions=output:$NEW_PORT
56
57     INFO "Successfully migrated from { ip: $OLD_IP, port: $OLD_PORT } to { ip: $NEW_IP, port: $NEW_PORT }"
58 done

```

Figure 4.28: ovs-ofctl migration

To sum up, there are no remote RYU controllers to manage the traffic, but the migration can be performed just adjusting flows.

4.4.1 Running the Scenario

Repeat the following commands two times, because two terminals are needed for the scenario.

```

$ cd SVMN/comnetsemu
$ vagrant ssh
$ cd comnetsmu/app/morphing_slices

```

Then, follow the steps:

- Generate the network from the topology.yaml file in the first terminal:

```
$ sudo python3 topology.py --file scenarios/2/topology.yaml
```

The output should be similar to figure 4.29 executing the Mininet command *dump*.

```

mininet> dump
<DockeHost m0: m0-eth0:10.0.0.10 pid=4066>
<DockeHost c0: c0-eth0:192.168.0.10 pid=4178>
<DockeHost s0: s0-eth0:192.168.0.100,s0-sw1-admin:None pid=4288>
<DockeHost s1: s1-eth0:192.168.0.100,s1-sw1-admin:None pid=4399>
<OVSSwitch sw0: lo:127.0.0.1,sw0-eth1:None,sw0-eth2:None,sw0-eth3:None pid=4011>
<OVSSwitch sw1: lo:127.0.0.1,sw1-eth1:None,sw1-s0-admin:None,sw1-s1-admin:None pid=4014>
<Controller controller0: 127.0.0.1:6653 pid=4003>

```

Figure 4.29: Mininet dump

- In the second terminal, run the ovs-ofctl script called *openflow.sh*.

```
$ scenarios/2/openflow.sh
```

```
vagrant@comnetsemu:~/comnetsemu/app/morphing_slices$ scenarios/2/openflow.sh
[2022-07-02 14:47:08.458][INFO ][main:27 ] Creating OpenFlow admin flow
[2022-07-02 14:47:08.498][INFO ][main:31 ] Creating OpenFlow data flow
Press 'Enter' to migrate or 'q' to exit
```

Figure 4.30: *openflow.sh* output

Any request can be performed to the enabled server in the terminal where the network topology has been deployed. So, the client can increment the counter.

```
$ c0 curl -X POST 192.168.0.100/api/counter
```

```
mininet> c0 curl -X POST 192.168.0.100/api/counter
{
  "counter": 1
}
mininet> c0 curl -X POST 192.168.0.100/api/counter
{
  "counter": 2
}
mininet> c0 curl -X POST 192.168.0.100/api/counter
{
  "counter": 3
}
```

Figure 4.31: Client increments counter

Flows can be checked in another terminal typing:

```
$ sudo ovs-ofctl dump-flows sw0
```

```
vagrant@comnetsemu:~/comnetsemu/app/morphing_slices$ sudo ovs-ofctl dump-flows sw0
cookie=0x0, duration=477.109s, table=0, n_packets=17, n_bytes=1365, in_port="sw0-eth1" actions=output:"sw0-eth2"
cookie=0x0, duration=477.095s, table=0, n_packets=17, n_bytes=1650, in_port="sw0-eth2" actions=output:"sw0-eth1"
cookie=0x0, duration=477.082s, table=0, n_packets=0, n_bytes=0, in_port="sw0-eth3" actions=output:"sw0-eth1"
```

Figure 4.32: Check flows

As shown in figure 4.32, there is no traffic in the flows with *in_port sw0-eth3*, because it is related to the server *s1* which is disabled at the beginning. The migration can be done just pressing *Enter* in the *ovs-ofctl* terminal.

```
vagrant@comnetsemu:~/comnetsemu/app/morphing_slices$ scenarios/2/openflow.sh
[2022-07-02 14:47:08.458][INFO ][main:27 ] Creating OpenFlow admin flow
[2022-07-02 14:47:08.498][INFO ][main:31 ] Creating OpenFlow data flow
Press 'Enter' to migrate or 'q' to exit
[2022-07-02 14:58:01.210][INFO ][main:49 ] Migrating from { ip: 10.0.0.100, port: 2 } to { ip: 10.0.0.101, port: 3 }
[2022-07-02 14:58:01.383][INFO ][main:57 ] Successfully migrated from { ip: 10.0.0.100, port: 2 } to { ip: 10.0.0.101, port: 3 }
Press 'Enter' to migrate or 'q' to exit
```

Figure 4.33: Perform migration

The client can continue to increment the counter where it has left without noticing the migration.

```
mininet> c0 curl -X POST 192.168.0.100/api/counter
{
  "counter": 1
}
mininet> c0 curl -X POST 192.168.0.100/api/counter
{
  "counter": 2
}
mininet> c0 curl -X POST 192.168.0.100/api/counter
{
  "counter": 3
}
mininet> c0 curl -X POST 192.168.0.100/api/counter
{
  "counter": 4
}
mininet> c0 curl -X POST 192.168.0.100/api/counter
{
  "counter": 5
}
mininet> c0 curl -X POST 192.168.0.100/api/counter
{
  "counter": 6
}
```

Figure 4.34: *c0* increments the counter after migration

Now, there is traffic in the flows with *in_port sw0-eth3*, because the server *s1* has become enabled.

```
vagrant@comnetsemu:~/comnetsemu/app/morphing_slices$ sudo ovs-ofctl dump-flows sw0
cookie=0x0, duration=893.809s, table=0, n_packets=17, n_bytes=1650, in_port="sw0-eth2" actions=output:"sw0-eth1"
cookie=0x0, duration=893.796s, table=0, n_packets=16, n_bytes=1608, in_port="sw0-eth3" actions=output:"sw0-eth1"
cookie=0x0, duration=893.823s, table=0, n_packets=33, n_bytes=2688, in_port="sw0-eth1" actions=output:"sw0-eth3"
```

Figure 4.35: Recheck flows

4.5 Scenario 3

The Scenario 3 is located in the folder *src/scenarios/3* and the network topology is reported in figure 4.36.

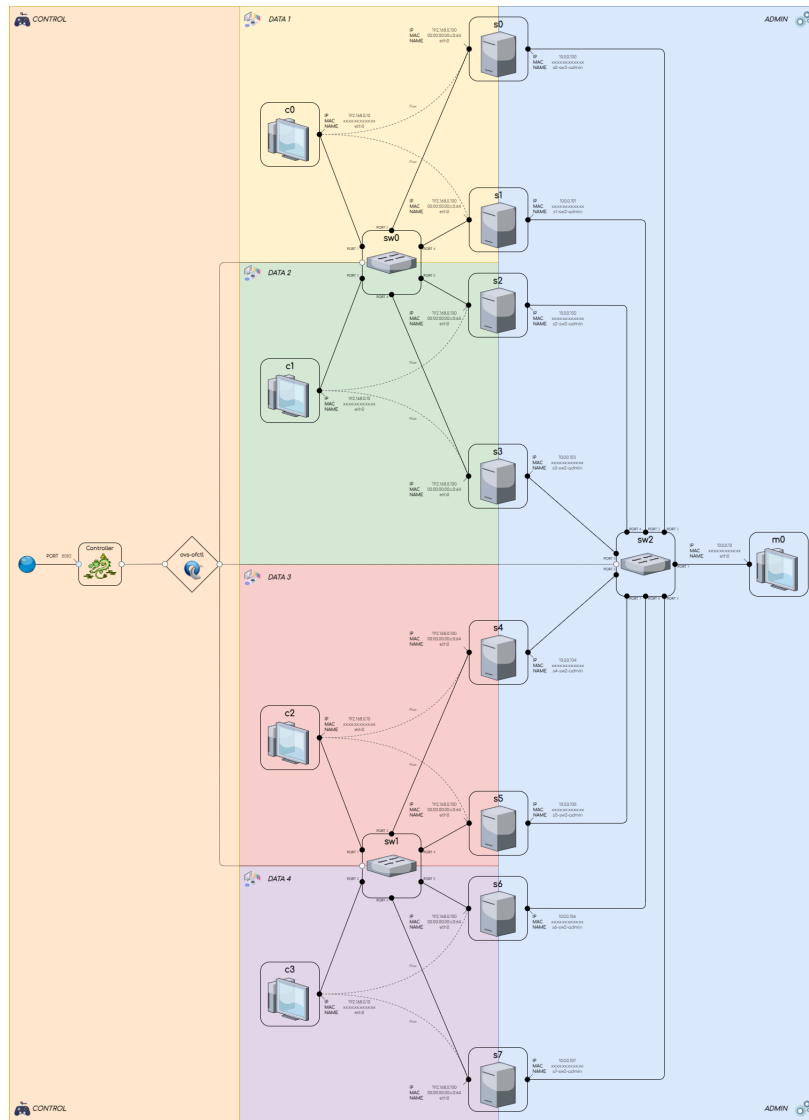


Figure 4.36: Network topology

There are a lot of devices in the network to prove the project scalability:

- Client *c0*, *c1*, *c2*, *c3*: they perform the HTTP requests to the enabled server in the correspondent slice.
- Servers *s0*, *s1*, *s2*, *s3*, *s4*, *s5*, *s6*, *s7*: they implement the Flask web application. By default, *s0*, *s2*, *s4*, *s6* are the enabled servers and *s1*, *s3*, *s5*, *s7* are disabled.
- Switch *sw0*: it connects *c0* with the two servers *s0* and *s1*. Moreover, it connects *c1* with the two servers *s2* and *s3*.
- Switch *sw1*: it connects *c2* with the two servers *s4* and *s5*. Moreover, it connects *c3* with the two servers *s6* and *s7*.
- Manager *m0*: it emulates the network administrator and executes the command for the migration.
- Switch *sw2*: it connects *m0* with the all the servers.
- Controller 0: it handles the network traffic.

The OpenFlow protocol is implemented by `ovs-ofctl` which defines slices and flows for the entire network. Slices and flows are defined by the script `src/scenarios/3/openflow.sh`.

```

99 # OpenFlow admin slice
100 INFO "Create admin slice on switch sw2"
101 sudo ovs-vsctl set port sw2-eth1 qos=@admin -- \
102 --id=@admin create qos type=linux-htb \
103 queues:777=@adm -- \
104 --id=@adm create queue other-config:max-rate=1000000
105 # OpenFlow admin flow
106 INFO "Creating OpenFlow admin flow"
107 ofctl_exec add-flow sw2 actions=set_queue:777,normal
108
109 # OpenFlow data slices sw0
110 INFO "Create slices data_1 and data_2 on switch sw0"
111 sudo ovs-vsctl set port sw0-eth1 qos=@data -- \
112 --id=@data create qos type=linux-htb \
113 queues:123=@data1 \
114 queues:234=@data2 -- \
115 --id=@data1 create queue other-config:max-rate=1000000 -- \
116 --id=@data2 create queue other-config:max-rate=1000000
117 # OpenFlow data_1 flows
118 INFO "Creating OpenFlow data_1 flows"
119 ofctl_exec add-flow sw0 in_port=1,actions=set_queue:123,output:1,${SLICE_1_SERVERS_PORT[SLICE_1_IDX_SERVER]}
120 ofctl_exec add-flow sw0 in_port=3,actions=set_queue:123,output:1
121 ofctl_exec add-flow sw0 in_port=4,actions=set_queue:123,output:1
122 # OpenFlow data_2 flows
123 INFO "Creating OpenFlow data_2 flows"
124 ofctl_exec add-flow sw0 in_port=2,actions=set_queue:234,output:2,${SLICE_2_SERVERS_PORT[SLICE_2_IDX_SERVER]}
125 ofctl_exec add-flow sw0 in_port=5,actions=set_queue:234,output:2
126 ofctl_exec add-flow sw0 in_port=6,actions=set_queue:234,output:2
127
128 # OpenFlow data slices sw1
129 INFO "Create slices data_3 and data_4 on switch sw1"
130 sudo ovs-vsctl set port sw1-eth1 qos=@data -- \
131 --id=@data create qos type=linux-htb \
132 queues:123=@data3 \
133 queues:234=@data4 -- \
134 --id=@data3 create queue other-config:max-rate=1000000 -- \
135 --id=@data4 create queue other-config:max-rate=1000000
136 # OpenFlow data_3 flows
137 INFO "Creating OpenFlow data_3 flows"
138 ofctl_exec add-flow sw1 in_port=1,actions=set_queue:123,output:1,${SLICE_3_SERVERS_PORT[SLICE_3_IDX_SERVER]}
139 ofctl_exec add-flow sw1 in_port=3,actions=set_queue:123,output:1
140 ofctl_exec add-flow sw1 in_port=4,actions=set_queue:123,output:1
141 # OpenFlow data_4 flows
142 INFO "Creating OpenFlow data_4 flows"
143 ofctl_exec add-flow sw1 in_port=2,actions=set_queue:234,output:2,${SLICE_4_SERVERS_PORT[SLICE_4_IDX_SERVER]}
144 ofctl_exec add-flow sw1 in_port=5,actions=set_queue:234,output:2
145 ofctl_exec add-flow sw1 in_port=6,actions=set_queue:234,output:2

```

Figure 4.37: `ovs-ofctl` slices and flows

There are the following slices:

- *data 1*, *data 2*, *data 3*, *data 4*: in each *data* slice, there are one client and two servers.
- *admin*: where the manager *m0* handles the server migration for every *data* slice.

When the script `src/scenarios/3/openflow.sh` runs, a while loop waits for the user input to perform the migration. The user can migrate in two modes:

- Update: the flow is updated with the new *in_port* attribute.
- Delete: the flow is removed and created with the new *in_port* attribute.

After the user picks the mode and in which slice to migrate:

- the manager *m0* performs the server migration via the Docker command line.
- the controller modifies the network traffic updating the flow, performing an HTTP post request to the *migrator.py* Flask web server described below.

```
113 # Docker manager
114 docker exec -d m0 curl -X POST -H "Content-Type:application/json" -d '{ "server": "${HTTP_HOST_IP}" "${NEW_IP}/api/admin/migrate
115 # Controller flow
116 curl -X POST -H "Content-Type:application/json" -d '{ "mode": "${MIGRATION_MODE}", "dpid": "${DPATH_ID}", "in_port": "${CLIENT_PORT}", "out_port": "${NEW_PORT}" }' "localhost:CONTROLLER_PORT/api/migrate"
```

Figure 4.38: Operations for migration

So, `ovs-ofctl` only creates slices and flows without updating them. The migration is entirely handled by the Controller 0 for any *data* slice dynamically adjusting the flow to always connect the client to the enabled server. The Controller 0 stores how to forward the traffic for switch *sw0* and *sw1* connecting the different *in_ports* in the dictionary *self.in_to_out*.

```
32 self.in_to_out = {
33     1: {1: 3, 2: 5, 3: 1, 4: 1, 5: 2, 6: 2},
34     2: {1: 3, 2: 5, 3: 1, 4: 1, 5: 2, 6: 2},
35 }
```

Figure 4.39: Port mapping for switch *sw0* and *sw1*

Furthermore, the Controller 0 runs a Flask web server `src/secenarios/3/migrate.py` in a new python thread to handle incoming HTTP request to update the flows and mapping. As soon as the controller is launched, the Flask web server shown in the figure 4.40 starts.

```
9 @app.route("/api/migrate", methods=["POST"])
10 def migrate():
11     """
12     Migrate service.
13     """
14
15     # Request body
16     body = request.get_json(force=True)
17     # Mode
18     mode = int(body["mode"])
19     # Switch dpid
20     dpid = int(body["dpid"])
21     # In port
22     in_port = int(body["in_port"])
23     # Out port
24     out_port = int(body["out_port"])
25
26     # Notify callback
27     if cb:
28         cb(mode, dpid, in_port, out_port)
29
30     return ("", 200)
31
32
33 def start(port: int, callback):
34     global cb
35     cb = callback
36     app.run(port=port)
```

Figure 4.40: *migrator.py*

The *migrator.py* accepts HTTP post request at */api/migrate* and call the method *migration_cb* to update or modify the flows and mapping. The mode, datapath id, MAC address and the new in port has to be passed in the body of the request. The mode parameter can be update or delete. Then, the method *migration_cb* shown in the figure 4.41 executes the user request.

```

50 def migration_cb(self, mode: int, dpid: int, in_port: int, out_port: int):
51     datapath: Datapath = self.get_datapath(dpid)
52
53     if mode == MIGRATION_MODE_UPDATE:
54         match = datapath.ofproto_parser.OFPMatch(in_port=in_port)
55         actions = [datapath.ofproto_parser.OFPActionOutput(out_port)]
56         self.update_flow(datapath, match, actions)
57     elif mode == MIGRATION_MODE_DELETE:
58         self.delete_flow(datapath, in_port)
59     else:
60         self.logger.warn(f"Unknown migration mode {mode}")
61
62     self.in_to_out[dpid][in_port] = out_port

```

Figure 4.41: *migrator_cb*

Hence, the Controller 0 is able to update the network traffic and flows in any *data* slice on its own. At the beginning, the Controller 0 can create flows considering the incoming OpenFlow traffic with the method *add_flow* shown in the figure 4.42.

```

150 def add_flow(
151     self, datapath: Datapath, priority: int, match, actions, buffer_id=None
152 ):
153     ofproto = datapath.ofproto
154     parser = datapath.ofproto_parser
155
156     self.logger.info(
157         f"Add flow: {{ dpid: {datapath.id}, priority: {priority}, match: {match}, actions: {actions} }}"
158     )
159
160     if buffer_id:
161         mod = parser.OFPFlowMod(
162             datapath=datapath,
163             match=match,
164             command=ofproto.OFPFC_ADD,
165             priority=priority,
166             flags=ofproto.OFPFF_SEND_FLOW_REM,
167             actions=actions,
168             buffer_id=buffer_id,
169         )
170     else:
171         mod = parser.OFPFlowMod(
172             datapath=datapath,
173             match=match,
174             command=ofproto.OFPFC_ADD,
175             priority=priority,
176             flags=ofproto.OFPFF_SEND_FLOW_REM,
177             actions=actions,
178         )
179
180     datapath.send_msg(mod)

```

Figure 4.42: Create flow

The Controller 0 can update a flow with the method *update_flow* reported in figure the 4.43.


```

182     def update_flow(self, datapath: Datapath, match, actions):
183         ofproto = datapath.ofproto
184         parser = datapath.ofproto_parser
185
186         self.logger.info(
187             f"Update flow: {{ dpid: {datapath.id}, match: {match}, actions: {actions} }}"
188         )
189
190         mod = parser.OFPFlowMod(
191             datapath=datapath,
192             match=match,
193             command=ofproto.OFPFC_MODIFY,
194             actions=actions,
195         )
196
197         datapath.send_msg(mod)

```

Figure 4.43: Update flow

The Controller 0 can delete a flow with the method *delete_flow* reported in figure the 4.44.

```

199     def delete_flow(self, datapath: Datapath, in_port: int):
200         ofproto = datapath.ofproto
201         parser = datapath.ofproto_parser
202
203         self.logger.info(
204             f"Delete flow: {{ dpid: {datapath.id}, in_port={in_port}, out_port: {self.in_to_out[datapath.id][in_port]} }}"
205         )
206
207         match = parser.OFPMatch(in_port=in_port)
208         mod = parser.OFPFlowMod(
209             datapath=datapath,
210             match=match,
211             command=ofproto.OFPFC_DELETE,
212         )
213
214         datapath.send_msg(mod)

```

Figure 4.44: Delete flow

After deleting a flow, it will be recreated again with new *in_port* attribute considering the incoming OpenFlow traffic.

4.5.1 Running the Scenario

Repeat the following commands three times, because three terminals in the Comnetsemu machine are needed.

```

$ cd SVMN/comnetsemu
$ vagrant ssh
$ cd comnetsmu/app/morphing_slices

```

Then, follow the steps:

1. Generate the network from the topology.yaml file in the first terminal:

```
$ sudo python3 topology.py --file scenarios/3/topology.yaml
```

The output should be similar to figure 4.45 executing the Mininet command *dump*.

```

mininet> dump
<DockerHost m0: m0-eth0:10.0.0.10 pid=5652>
<DockerHost c0: c0-eth0:192.168.0.10 pid=5767>
<DockerHost c1: c1-eth0:192.168.0.10 pid=5881>
<DockerHost c2: c2-eth0:192.168.0.10 pid=5988>
<DockerHost c3: c3-eth0:192.168.0.10 pid=6095>
<DockerHost s0: s0-eth0:192.168.0.100, s0-sw2-admin:None pid=6195>
<DockerHost s1: s1-eth0:192.168.0.100, s1-sw2-admin:None pid=6299>
<DockerHost s2: s2-eth0:192.168.0.100, s2-sw2-admin:None pid=6400>
<DockerHost s3: s3-eth0:192.168.0.100, s3-sw2-admin:None pid=6506>
<DockerHost s4: s4-eth0:192.168.0.100, s4-sw2-admin:None pid=6610>
<DockerHost s5: s5-eth0:192.168.0.100, s5-sw2-admin:None pid=6707>
<DockerHost s6: s6-eth0:192.168.0.100, s6-sw2-admin:None pid=6809>
<DockerHost s7: s7-eth0:192.168.0.100, s7-sw2-admin:None pid=6910>
<OVSSwitch sw0: lo:127.0.0.1, sw0-eth1:None, sw0-eth2:None, sw0-eth3:None, sw0-eth4:None, sw0-eth5:None, sw0-eth6:None pid=5592>
<OVSSwitch sw1: lo:127.0.0.1, sw1-eth1:None, sw1-eth2:None, sw1-eth3:None, sw1-eth4:None, sw1-eth5:None, sw1-eth6:None pid=5595>
<OVSSwitch sw2: lo:127.0.0.1, sw2-eth1:None, sw2-s0-admin:None, sw2-s1-admin:None, sw2-s2-admin:None, sw2-s3-admin:None, sw2-s4-admin:None, sw2-s5-admin:None, sw2-s6-admin:None, sw2-s7-admin:None pid=5598>
<RemoteController controller0: 127.0.0.1:6653 pid=5585>

```

Figure 4.45: Mininet dump

2. In the second terminal, run the ovs-ofctl script called *openflow.sh*.

```
$ scenarios/3/openflow.sh
```

Pick between *UPDATE* or *DELETE* mode for the migration and the output should be similar to figure 4.46.

```
vagrant@comnetsemu:~/comnetsemu/app/morphing_slices$ scenarios/3/openflow.sh
Select migration mode [1->UPDATE|2->DELETE]: 1
[2022-07-03 09:32:25.763][INFO ][main:100] Create admin slice on switch sw2
d61bbbec-551a-4e83-bcf8-3ac07426819f
1c0ec60f-e45e-4d95-a715-7836f5e9caa1
[2022-07-03 09:32:25.801][INFO ][main:106] Creating OpenFlow admin flow
[2022-07-03 09:32:25.831][INFO ][main:110] Create slices data_1 and data_2 on switch sw0
a9796a96-8c64-48ca-9a26-e3250473e766
e0b97929-5224-43a9-bf7a-834bc01a24e4
af7d8606-d410-4984-bcf8-c5a3b9e455ed
[2022-07-03 09:32:25.859][INFO ][main:118] Creating OpenFlow data_1 flows
[2022-07-03 09:32:25.923][INFO ][main:123] Creating OpenFlow data_2 flows
[2022-07-03 09:32:25.968][INFO ][main:129] Create slices data_3 and data_4 on switch sw1
ec14fcb6-4ced-4563-a8a1-b1a644efd70a
2805e177-f275-4805-83b2-9d861de81ac4
90ece96f-e741-46bd-8abb-e7e0345b51a9
[2022-07-03 09:32:25.990][INFO ][main:137] Creating OpenFlow data_1 flows
[2022-07-03 09:32:26.039][INFO ][main:142] Creating OpenFlow data_2 flows
Press 'Enter' to migrate or 'q' to exit|
```

Figure 4.46: Run *openflow.sh*

3. In the third terminal, run the RYU controller:

```
$ scripts/ryu.sh --controller scenarios/3/controller.py --port 8082 \
--config scenarios/3/controller.cfg
```

```
vagrant@comnetsemu:~/comnetsemu/app/morphing_slices$ scripts/ryu.sh --controller scenarios/3/controller.py --port 8082 --config scenarios/3/controller.cfg
[2022-07-03 09:38:06.816][INFO ][main:74] Running ryu: { controller: scenarios/3/controller.py, ofport: 8082, config: scenarios/3/controller.cfg } loading app /u
sr/local/lib/python3.8/dist-packages/ryu/app/gui_topology/gui_topology.py
loading app scenarios/3/controller.py
loading app ryu.app.rest_topology
loading app ryu.app.ofctl_rest
loading app ryu.app.ws_topology
loading app ryu.controller.ofp_handler
creating context wsgi
instantiating app None of DPSet
creating context dpset
creating context switches
instantiating app /usr/local/lib/python3.8/dist-packages/ryu/app/gui_topology/gui_topology.py of GUIServiceApp
instantiating app scenarios/3/controller.py of Controller
* Serving Flask app 'migrator' (lazy loading)
* Environment: production
WARNING: This is a development server. Do not use it in a production deployment.
Use a production WSGI server instead.
* Debug mode: off
* Running on http://127.0.0.1:8082 (Press CTRL+C to quit)
instantiating app ryu.app.rest_topology of TopologyAPI
instantiating app ryu.app.ofctl_rest of RestStatsd
instantiating app ryu.app.ws_topology of WebSocketTopology
instantiating app ryu.controller.ofp_handler of OFPHandler
(8084) wsgi starting up on http://0.0.0.0:8082
OPFSwitchFeatures: { dpid: 3, buffers: 0, tables: 254, capabilities: 199, ports: [{ port: 65534, name: sw2, hw_addr: e6:20:21:91:dd:4a }, { port: 8, name: sw2-s6-admin, h
w_addr: aa:f9:e1:85:14:1f }, { port: 5, name: sw2-s3-admin, hw_addr: 92:c8:b1:e7:68:6d }, { port: 9, name: sw2-s7-admin, hw_addr: a2:a8:0b:4b:94:bd }, { port: 4, name: sw
2-s2-admin, hw_addr: 06:08:4c:c9:a3:09 }, { port: 1, name: sw2-eth1, hw_addr: 72:57:47:bd:fe:b5 }, { port: 7, name: sw2-s5-admin, hw_addr: ce:c6:78:cd:d3:51 }, { port: 6,
name: sw2-s4-admin, hw_addr: 9a:47:68:fa:3b:5b }, { port: 2, name: sw2-s8-admin, hw_addr: a6:56:c8:98:89:c4 }, { port: 3, name: sw2-s1-admin, hw_addr: 2a:14:5e:09:56:ca
}] }
OPFSwitchFeatures: { dpid: 1, buffers: 0, tables: 254, capabilities: 199, ports: [{ port: 65534, name: sw0, hw_addr: 66:23:1b:9d:a9:4b }, { port: 4, name: sw0-eth4, hw_ad
dr: 6a:7a:36:2f:07:03 }, { port: 6, name: sw0-eth6, hw_addr: 02:77:38:9b:c3:55 }, { port: 1, name: sw0-eth3, hw_addr: 26:4e:d5:1a:22:a3 }, { port: 2, name: sw0-eth2, hw_a
ddr: 4e:5f:2f:07:71:86 }, { port: 5, name: sw0-eth5, hw_addr: 3a:55:d1:1f:c6:5d }, { port: 3, name: sw0-eth3, hw_addr: 22:46:01:a2:dc:08 }] }
SwitchEnter: { dpid: 3 }
SwitchEnter: { dpid: 1 }
OPFSwitchFeatures: { dpid: 2, buffers: 0, tables: 254, capabilities: 199, ports: [{ port: 65534, name: sw1, hw_addr: d6:68:d8:ec:d8:4a }, { port: 6, name: sw1-eth6, hw_a
ddr: f6:04:c9:ae:0a:81 }, { port: 4, name: sw1-eth4, hw_addr: de:18:44:0c:e1:80 }, { port: 1, name: sw1-eth1, hw_addr: 66:01:49:4e:23:88 }, { port: 2, name: sw1-eth2, hw_a
ddr: 6e:e4:b8:ce:d1:37 }, { port: 5, name: sw1-eth5, hw_addr: c2:4f:a7:74:97:b3 }, { port: 3, name: sw1-eth3, hw_addr: 1a:f7:98:e4:e4:75 }] }
SwitchEnter: { dpid: 2 }
```

Figure 4.47: Run RYU Controller 0

The RYU Graphical User Interface can be accessed from the browser at the link <http://localhost:8082>.

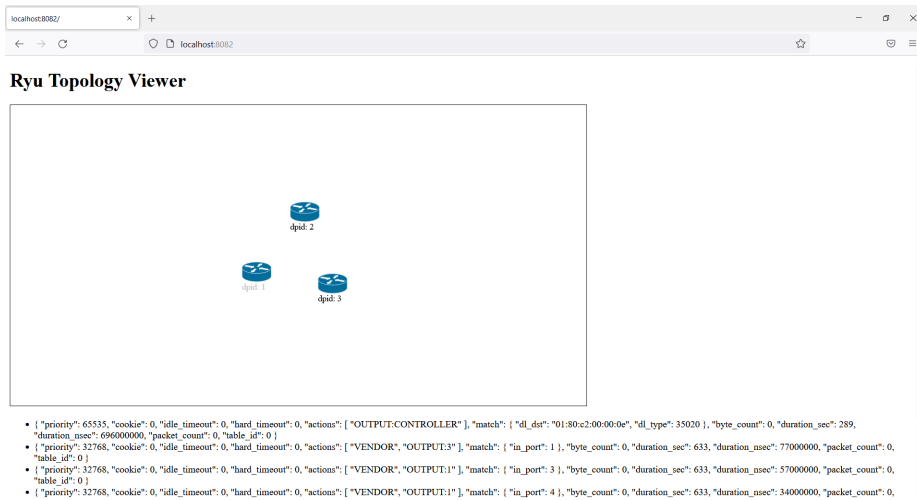


Figure 4.48: Run RYU at textithttp://localhost:8082

Requests can be performed by client *c0*, *c1*, *c2*, *c3*.

```
$ c1 curl -X POST 192.168.0.100/api/counter
```

```
mininet> c1 curl -X POST 192.168.0.100/api/counter
{"counter": 1}
mininet> c1 curl -X POST 192.168.0.100/api/counter
{"counter": 2}
```

Figure 4.49: Client *c1* increments counter

Slices and flows can be checked in another terminal typing:

```
$ sudo ovs-ofctl dump-flows sw0
$ sudo ovs-ofctl dump-flows sw1
```

```
vagrant@connetsemu:~/connetsemu/app/morphing_slices$ sudo ovs-ofctl dump-flows sw0
cookie=0x0, duration=983.589s, table=0, n_packets=0, n_bytes=0, priority=65535,dl_dst=01:80:c2:00:00:0e,dl_type=0x88cc actions=CONTROLLER:60
cookie=0x0, duration=1326.970s, table=0, n_packets=19, n_bytes=1521, in_port="sw0-eth1" actions=set_queue:123,output:"sw0-eth3"
cookie=0x0, duration=1326.958s, table=0, n_packets=19, n_bytes=1806, in_port="sw0-eth3" actions=set_queue:123,output:"sw0-eth1"
cookie=0x0, duration=1326.927s, table=0, n_packets=0, n_bytes=0, in_port="sw0-eth4" actions=set_queue:123,output:"sw0-eth1"
cookie=0x0, duration=1326.910s, table=0, n_packets=12, n_bytes=962, in_port="sw0-eth2" actions=set_queue:234,output:"sw0-eth5"
cookie=0x0, duration=1326.897s, table=0, n_packets=12, n_bytes=1152, in_port="sw0-eth5" actions=set_queue:234,output:"sw0-eth2"
cookie=0x0, duration=1326.883s, table=0, n_packets=0, n_bytes=0, in_port="sw0-eth6" actions=set_queue:234,output:"sw0-eth2"
vagrant@connetsemu:~/connetsemu/app/morphing_slices$ sudo ovs-ofctl dump-flows sw1
cookie=0x0, duration=985.203s, table=0, n_packets=0, n_bytes=0, priority=65535,dl_dst=01:80:c2:00:00:0e,dl_type=0x88cc actions=CONTROLLER:60
cookie=0x0, duration=1328.465s, table=0, n_packets=0, n_bytes=0, in_port="sw1-eth1" actions=set_queue:123,output:"sw1-eth3"
cookie=0x0, duration=1328.451s, table=0, n_packets=0, n_bytes=0, in_port="sw1-eth3" actions=set_queue:123,output:"sw1-eth1"
cookie=0x0, duration=1328.433s, table=0, n_packets=0, n_bytes=0, in_port="sw1-eth4" actions=set_queue:123,output:"sw1-eth1"
cookie=0x0, duration=1328.416s, table=0, n_packets=0, n_bytes=0, in_port="sw1-eth2" actions=set_queue:234,output:"sw1-eth5"
cookie=0x0, duration=1328.402s, table=0, n_packets=0, n_bytes=0, in_port="sw1-eth5" actions=set_queue:234,output:"sw1-eth2"
cookie=0x0, duration=1328.388s, table=0, n_packets=0, n_bytes=0, in_port="sw1-eth6" actions=set_queue:234,output:"sw1-eth2"
```

Figure 4.50: Check slices and flows

In the second terminal, the migration can be performed picking the *data* slice corresponding to the client. In the example, the *data* slice of the client *c1* is the number 2 according to the network topology.

```
Press 'Enter' to migrate or 'q' to exit
Select slice to migrate [1|2|3|4]: 2
[2022-07-03 09:58:03.875][INFO] [main:211] Migrating from { ip: 10.0.0.102, port: 5 } to { ip: 10.0.0.103, port: 6 }
[2022-07-03 09:58:04.111][INFO] [main:227] Successfully migrated slice 2 (data_2) from { ip: 10.0.0.102, port: 5 } to { ip: 10.0.0.103, port: 6 }
```

Figure 4.51: Perform migration

So, the flow is updated by the Controller 0 with the method *update_flow* shown in the figure 4.43. The client *c1* can continue to increment the counter where it has left without noticing the migration.

```
mininet> c1 curl -X POST 192.168.0.100/api/counter
{
  "counter": 1
}
mininet> c1 curl -X POST 192.168.0.100/api/counter
{
  "counter": 2
}
mininet> c1 curl -X POST 192.168.0.100/api/counter
{
  "counter": 3
}
mininet> c1 curl -X POST 192.168.0.100/api/counter
{
  "counter": 4
}
mininet> c1 curl -X POST 192.168.0.100/api/counter
{
  "counter": 5
}
```

Figure 4.52: *c1* increments counter after migration

The update operation can be checked in the Controller 0 logs as shown in the figure 4.53.

```
10.0.2.2 - - [03/Jul/2022 09:42:43] "GET / HTTP/1.1" 200 177 0.000695
10.0.2.2 - - [03/Jul/2022 09:42:43] "GET /ryu.topology.css HTTP/1.1" 200 514 0.003409
(8014) accepted ("10.0.2.2", 52495)
10.0.2.2 - - [03/Jul/2022 09:42:43] "GET /ryu.topology.js HTTP/1.1" 200 8652 0.005576
(8014) accepted ("10.0.2.2", 52502)
10.0.2.2 - - [03/Jul/2022 09:42:43] "GET /v1.0/topology/switches HTTP/1.1" 200 2494 0.003226
10.0.2.2 - - [03/Jul/2022 09:42:43] "GET /v1.0/topology/links HTTP/1.1" 200 133 0.000701
10.0.2.2 - - [03/Jul/2022 09:42:43] "GET /router.svg HTTP/1.1" 204 177 0.001304
10.0.2.2 - - [03/Jul/2022 09:42:46] "GET /stats/flow/1 HTTP/1.1" 200 1799 0.004432
10.0.2.2 - - [03/Jul/2022 09:42:49] "GET /stats/flow/3 HTTP/1.1" 200 631 0.003309
10.0.2.2 - - [03/Jul/2022 09:42:50] "GET /stats/flow/2 HTTP/1.1" 200 1795 0.004661
10.0.2.2 - - [03/Jul/2022 09:42:50] "GET /stats/flow/1 HTTP/1.1" 200 1795 0.011707
10.0.2.2 - - [03/Jul/2022 09:42:56] "GET /v1.0/topology/ms HTTP/1.1" 200 0 12.970134
10.0.2.2 - - [03/Jul/2022 09:42:56] "GET / HTTP/1.1" 304 177 0.001197
(8014) accepted ("10.0.2.2", 52505)
10.0.2.2 - - [03/Jul/2022 09:42:56] "GET /v1.0/topology/switches HTTP/1.1" 200 2494 0.010398
10.0.2.2 - - [03/Jul/2022 09:42:56] "GET /v1.0/topology/links HTTP/1.1" 200 133 0.001647
10.0.2.2 - - [03/Jul/2022 09:42:58] "GET /stats/flow/1 HTTP/1.1" 200 1793 0.004632
10.0.2.2 - - [03/Jul/2022 09:43:20] "GET /v1.0/topology/ms HTTP/1.1" 200 0 23.485612
Update Flow: { 'id': 1, 'match': OFMatch(dl_dst=b'\x00\x00\x00\x00\x00', dl_src=b'\x00\x00\x00\x00\x00', dl_type=0, dl_vlan=0, dl_vlan_pcp=0, in_port=2, nw_dst=0, nw_proto=0, nw_src=0, nw_tos=0, tp_dst=0, tp_src=0, wildcard=194392), 'actions': [OFActionOutput(max_len=65509, port=3)] }
127.0.0.1 - - [03/Jul/2022 09:58:04] "POST /api/migrate HTTP/1.1" 200
```

Figure 4.53: Controller 0 logs

Then, slices and flows can be rechecked again.

```
$ sudo ovs-ofctl dump-flows sw0
$ sudo ovs-ofctl dump-flows sw1
```

```
vagrant@connetsemu:~/connetsemu/app/morphing_slices$ sudo ovs-ofctl dump-flows sw0
cookie=0x0, duration=1600.787s, table=0, n_packets=0, n_bytes=0, priority=65535,dl_dst=01:80:c2:00:00:0e,dl_type=0x88cc actions=CONTROLLER:60
cookie=0x0, duration=1944.168s, table=0, n_packets=19, n_bytes=1521, in_port="sw0-eth1" actions=set_queue:123,output:"sw0-eth3"
cookie=0x0, duration=1944.148s, table=0, n_packets=19, n_bytes=1806, in_port="sw0-eth3" actions=set_queue:123,output:"sw0-eth1"
cookie=0x0, duration=1944.125s, table=0, n_packets=0, n_bytes=0, in_port="sw0-eth4" actions=set_queue:123,output:"sw0-eth1"
cookie=0x0, duration=1944.095s, table=0, n_packets=12, n_bytes=1152, in_port="sw0-eth5" actions=set_queue:234,output:"sw0-eth2"
cookie=0x0, duration=1944.081s, table=0, n_packets=18, n_bytes=1740, in_port="sw0-eth6" actions=set_queue:234,output:"sw0-eth2"
cookie=0x0, duration=1944.108s, table=0, n_packets=32, n_bytes=2549, in_port="sw0-eth2" actions=output:"sw0-eth6"
vagrant@connetsemu:~/connetsemu/app/morphing_slices$ sudo ovs-ofctl dump-flows sw1
cookie=0x0, duration=1601.841s, table=0, n_packets=0, n_bytes=0, priority=65535,dl_dst=01:80:c2:00:00:0e,dl_type=0x88cc actions=CONTROLLER:60
cookie=0x0, duration=1945.103s, table=0, n_packets=0, n_bytes=0, in_port="sw1-eth1" actions=set_queue:123,output:"sw1-eth3"
cookie=0x0, duration=1945.089s, table=0, n_packets=0, n_bytes=0, in_port="sw1-eth3" actions=set_queue:123,output:"sw1-eth1"
cookie=0x0, duration=1945.071s, table=0, n_packets=0, n_bytes=0, in_port="sw1-eth4" actions=set_queue:123,output:"sw1-eth1"
cookie=0x0, duration=1945.054s, table=0, n_packets=0, n_bytes=0, in_port="sw1-eth2" actions=set_queue:234,output:"sw1-eth5"
cookie=0x0, duration=1945.040s, table=0, n_packets=0, n_bytes=0, in_port="sw1-eth5" actions=set_queue:234,output:"sw1-eth2"
cookie=0x0, duration=1945.026s, table=0, n_packets=0, n_bytes=0, in_port="sw1-eth6" actions=set_queue:234,output:"sw1-eth2"
```

Figure 4.54: Recheck slices and flows

Chapter 5

Known Issues

The known issues faced during the project development are listed below:

- If something went wrong running the installation script *scripts/init.sh* on Windows OS as described in the chapter 1, the Virtual Box window should be kept in evidence by the mouse.
- The files for the RYU Graphical User Interface miss, installing RYU from pip. So, they are downloaded manually as shown in the figure 1.8 in the chapter 1.
- FlowVisor only implements the OpenFlow protocol version 1.0, which is the oldest one. Moreover, it does not allow to define directly the OpenFlow resulting action. It needs to be restarted to update any slice or flow and the built-in update command does not work as expected.
- RYU and Mininet documentation are not always clear. Indeed, it has been put a lot of effort to deploy the last RYU controller described in the scenario 3 at 4.5 in the chapter 4.

Conclusions

The project has been developed to be cross platform and run on any operating system. On Windows might be longer the process, but it is just a matter of downloading a plugin and configuring correctly the environment.

Then, a new development workflow has been proposed in the Comnetsemu environment. It is very interesting, because it allows you to develop the code and immediately run it in the Comnetsemu virtual machine. So, the developer does not need to take care of copying and pasting every time the code inside the Comnetsemu virtual machine.

The core of the project is the new topology building system described in the chapter 3. It is the innovation introduced by the project that can help to create Mininet networks in Comnetsemu easier. Furthermore, the network topology files are more readable and it is simple to update and debug them.

At the end, different scenarios have been tested and results have been reported. The first scenario is the most complicated in terms of service migration and OpenFlow protocol, because FlowVisor is deployed by a new Docker container. It might be difficult to develop it, because it needs to create the communication between the FlowVisor Docker container and the RYU controller. Instead, the third scenario is the best option in terms of service migration, because it only relies on the RYU controller to perform the service migration. Moreover, slices and flows are handled by ovs-ofctl and there is no need to deploy a new Docker container outside the Mininet network as for FlowVisor. So, the RYU controller can modify directly slices and flows without communicating with a Docker container in a new virtualization level. The second scenario is the most simple one. It is shown to introduce ovs-ofctl to lead at the final best scenario 3.

Bibliography

- [1] *CD Workflow*. URL: <https://docs.github.com/en/actions/deployment/about-deployments/about-continuous-deployment>.
- [2] *CI Workflow*. URL: <https://docs.github.com/en/actions/automating-builds-and-tests/about-continuous-integration>.
- [3] *Comnetsemu Repository*. URL: <https://git.comnets.net/public-repo/comnetsemu>.
- [4] *Containernet*. URL: <https://containernet.github.io/>.
- [5] *Dependabot*. URL: <https://docs.github.com/en/code-security/dependabot>.
- [6] *Docker*. URL: <https://www.docker.com/>.
- [7] *Flask*. URL: <https://flask.palletsprojects.com/en/2.1.x>.
- [8] *FlowVisor*. URL: <https://github.com/OPENNETWORKINGLAB/flowvisor>.
- [9] *GitHub Workflows*. URL: <https://docs.github.com/en/github-ae@latest/actions/using-workflows/about-workflows>.
- [10] *inotifywait*. URL: <https://linux.die.net/man/1/inotifywait>.
- [11] *Kubernetes*. URL: <https://kubernetes.io/>.
- [12] *Mininet*. URL: <http://mininet.org/>.
- [13] *OpenFlow Protocol*. URL: <https://en.wikipedia.org/wiki/OpenFlow>.
- [14] *ovs-ofctl*. URL: <https://www.man7.org/linux/man-pages/man8/ovs-ofctl.8.html>.
- [15] *Pre-commit*. URL: <https://pre-commit.com/>.
- [16] *PySerde*. URL: <https://github.com/yukinarit/pyserde>.
- [17] *rsync*. URL: <https://linux.die.net/man/1/rsync>.
- [18] *RYU Controller Documentation*. URL: <https://ryu.readthedocs.io/en/latest/>.

List of Figures

1	Introduction summary	2
1.1	First <i>init.sh</i> run	3
1.2	Packages required	4
1.3	Set Python virtual environment	4
1.4	Run and configure Comnetsemu	4
1.5	Update and install Comnetsemu packages	4
1.6	Build Docker images	5
1.7	Install Python packages in Comnetsemu	5
1.8	Fix RYU Python package	5
2.1	CI steps	7
2.2	CI result	8
2.3	CD steps	8
2.4	Folder generation	8
2.5	CD result	9
2.6	CD release	9
2.7	CD release download	9
2.8	Pre-commit YAML file	10
2.9	Dependabot YAML file	10
2.10	Development schema	10
2.11	Check Comnetsemu staus	11
2.12	Generate requirements	11
2.13	Rsync synchronization	11
2.14	Initialize Comnetsemu	11
2.15	Start watcher	11
2.16	Run <i>dev.sh</i>	12
2.17	Change occurrence	12
3.1	Main <i>topology.py</i>	13
3.2	Get and build topology	14
3.3	Start network and Mininet CLI	14
3.4	<i>getTopology</i> function	14
3.5	<i>fileFormat</i> function	14
3.6	<i>parseTopology</i> function	15
3.7	Topology object	15
3.8	Network object	15
3.9	Controller objects	15

3.10	Switch object	16
3.11	Network link object	16
3.12	Host object	16
3.13	Container object	16
3.14	Network interface object	16
3.15	<i>buildTopology</i> function	17
3.16	<i>buildControllers</i> function	18
3.17	<i>buildSwitches</i> function	18
3.18	<i>buildHosts</i> function	18
3.19	<i>buildNetworkLinks</i> function	19
3.20	<i>buildNetworkInterfaces</i> function	19
3.21	<i>buildContainers</i> function	19
3.22	<i>cleanTopology</i> function	20
3.23	Building topology output	20
4.1	Flask counter	21
4.2	Flask <i>/api/counter</i> [GET]	22
4.3	Flask <i>/api/counter</i> [POST]	22
4.4	Flask <i>/api/admin/disable</i> [POST]	22
4.5	Flask <i>/api/admin/migrate</i> [POST]	23
4.6	Flask middleware	23
4.7	Network topology	24
4.8	Slices and flows	25
4.9	Flask <i>flowvisor.py</i>	25
4.10	Run Flask <i>flowvisor.py</i>	25
4.11	Self-learning controller	26
4.12	Spawn <i>migrator.py</i>	26
4.13	<i>migrator.py</i> API	27
4.14	While loop migration	27
4.15	Mininet dump	28
4.16	Run FlowVisor container	28
4.17	Run <i>flowvisor.sh</i> migration script	28
4.18	Run controllers	29
4.19	RYU GUI at <i>http://localhost:8082</i>	29
4.20	RYU GUI at <i>http://localhost:8083</i>	30
4.21	Client increments counter	30
4.22	<i>s0</i> replies to client	30
4.23	FlowVisor migration	31
4.24	<i>c0</i> increments the counter after migration	31
4.25	<i>s1</i> replies to client	31
4.26	Network topology	32
4.27	ovs-ofctl flows	32
4.28	ovs-ofctl migration	33
4.29	Mininet dump	33
4.30	<i>openflow.sh</i> output	34
4.31	Client increments counter	34
4.32	Check flows	34
4.33	Perform migration	34
4.34	<i>c0</i> increments the counter after migration	34

4.35	Recheck flows	35
4.36	Network topology	35
4.37	ovs-ofctl slices and flows	36
4.38	Operations for migration	37
4.39	Port mapping for switch <i>sw0</i> and <i>sw1</i>	37
4.40	<i>migrator.py</i>	37
4.41	<i>migrator_cb</i>	38
4.42	Create flow	38
4.43	Update flow	39
4.44	Delete flow	39
4.45	Mininet dump	39
4.46	Run <i>openflow.sh</i>	40
4.47	Run RYU Controller 0	40
4.48	Run RYU at textithttp://localhost:8082	41
4.49	Client <i>c1</i> increments counter	41
4.50	Check slices and flows	41
4.51	Perform migration	41
4.52	<i>c1</i> increments counter after migration	42
4.53	Controller 0 logs	42
4.54	Recheck slices and flows	42